

A Practical Introduction to (Explainable) Graph Learning

Yong-Min Shin

School of Mathematics and Computing (Computational Science and Engineering)

Yonsei University

2025.04.02



수학계산학부(계산과학공학)

School of Mathematics and Computing
(Computational Science and Engineering)



이화여자대학교

EWHA WOMANS UNIVERSITY

Objectives

Part 1: A practical introduction to graphs and graph neural networks

1. Understanding of **graphs** as a **general data type**
2. Understanding of the general framework of **graph neural networks (GNNs)**

Part 2: Towards explainable graph learning with attention

1. Understanding the basic concepts of **explainable AI**
2. **Answer to the question: Can we understand graph attention networks using attention?**



Part 1: A practical introduction to graphs and graph neural networks

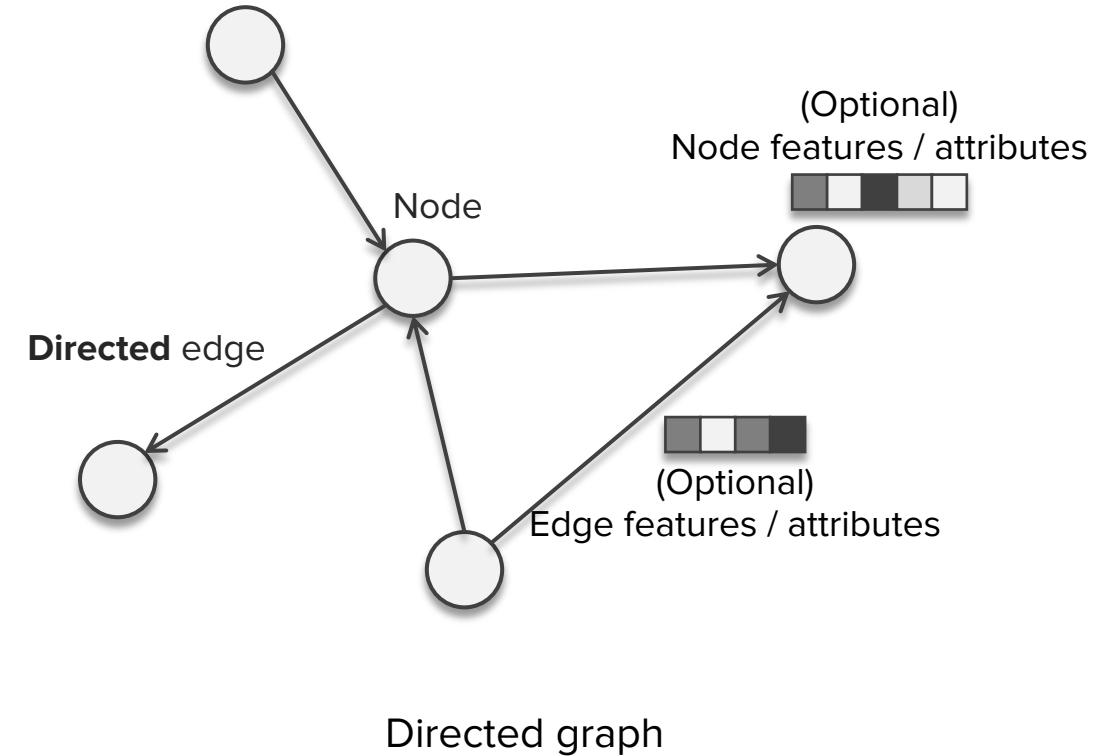
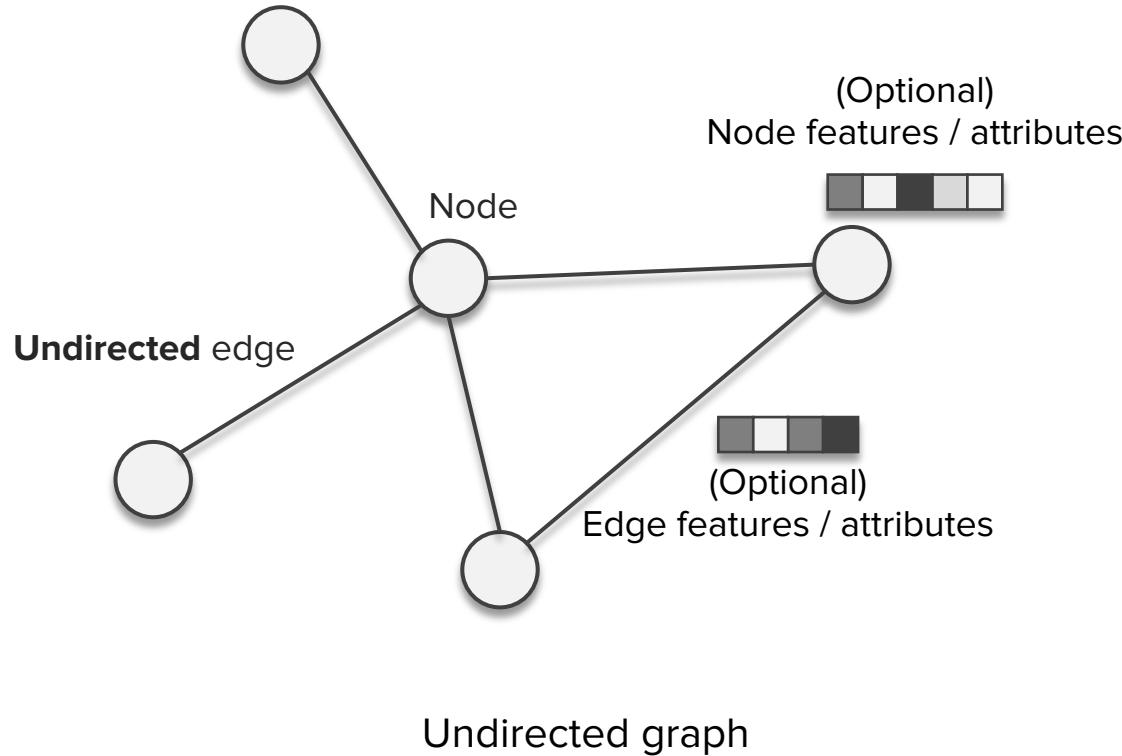
Understanding of graphs as a general data type

*This part is heavily influenced by one of my academic heros, Petar Veličković. These are some materials from his public materials that I have referred to:

- (Slide) Everything is Connected: Graph Neural Networks from the Ground Up (2021)
- (Blog) Graph & Geometric ML in 2024: Where We Are and What's Next (Part II – Applications)

Graphs as an abstract datatype

Graphs are an abstract type of data where nodes (entities) are connected by edges (connections)

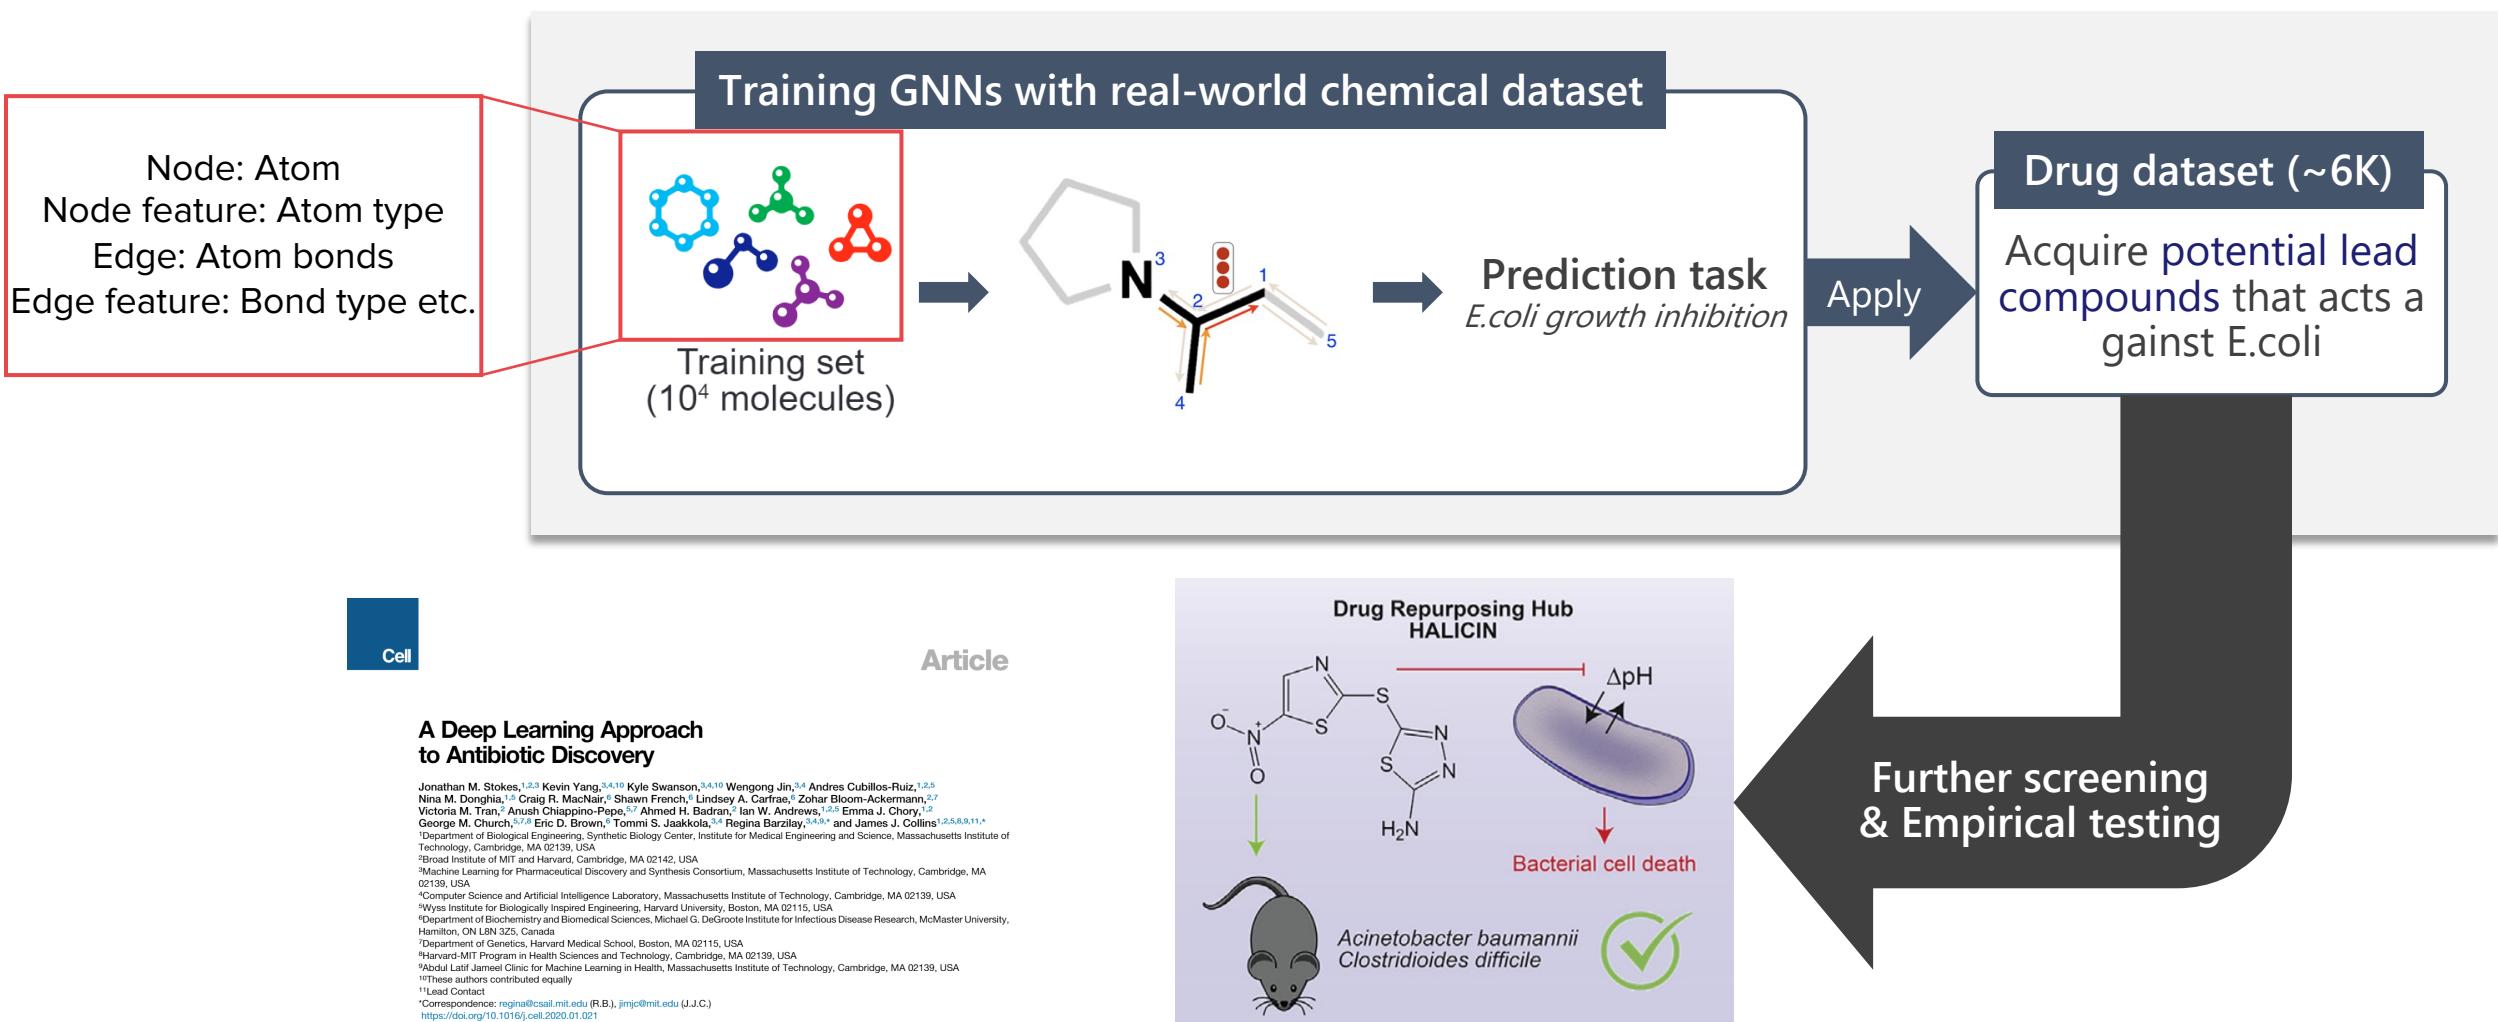


...But honestly, looking at this does not result in a **practical** understanding of graphs.

Therefore, we will look at **various applications** in the field of **graph machine learning** before moving our discussion further.

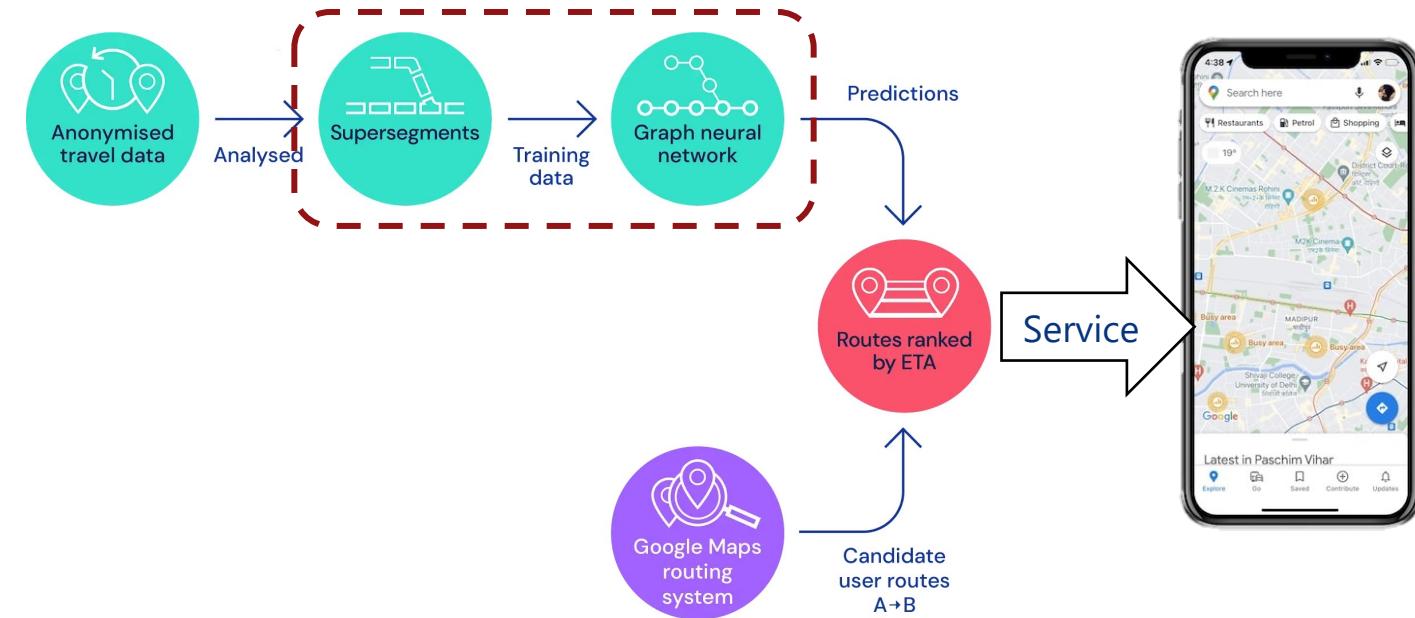
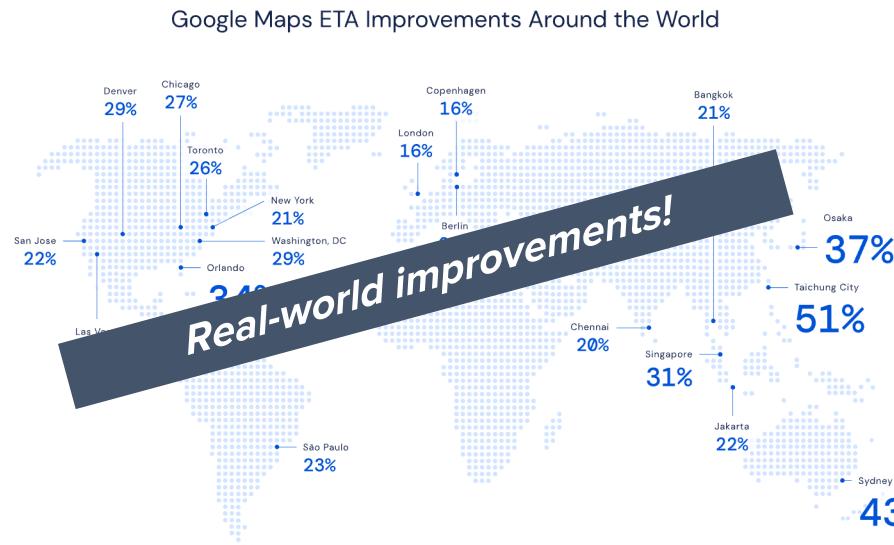
Area 1) Biology & Chemistry Research

Example 1: The discovery of Halicin, GNN-guided antibiotic discovery

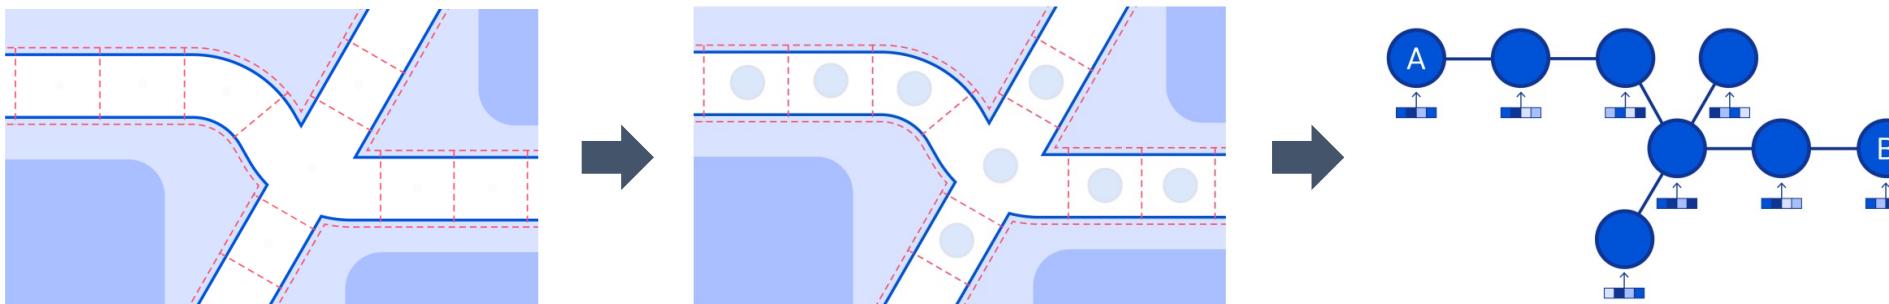


Area 2) ETA prediction

Example 2: DeepMind's improvement of Google map's ETA (Estimated Time of Arrival) prediction



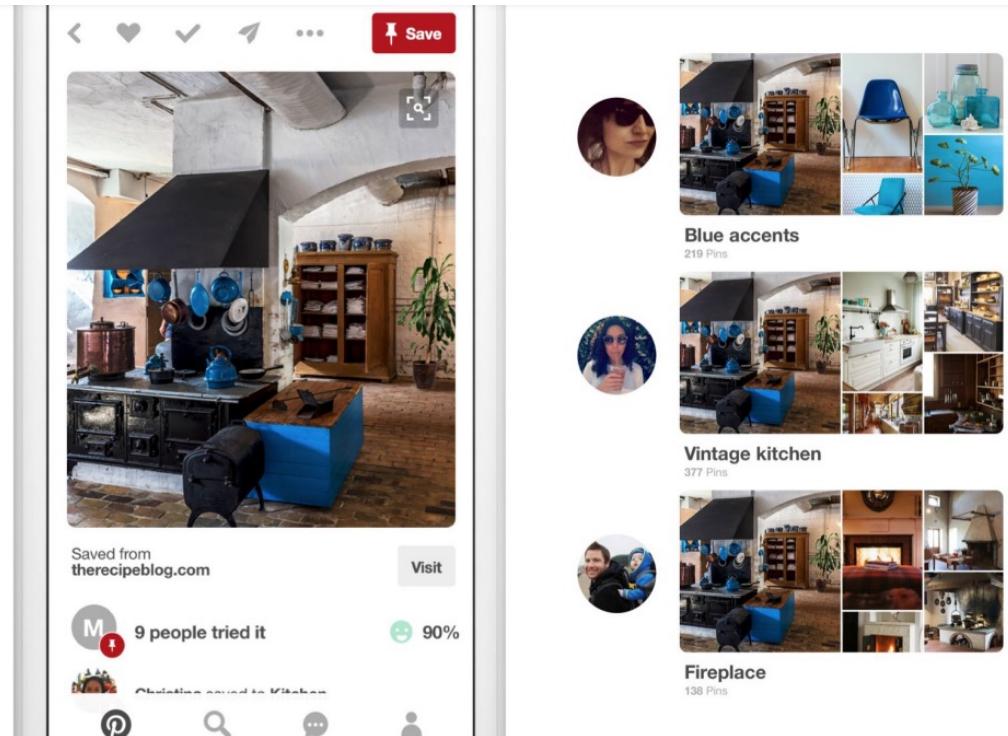
Unlike chemical datasets, constructing a graph is less straightforward.
In these cases, **how to construct the graph** is also a crucial contribution.



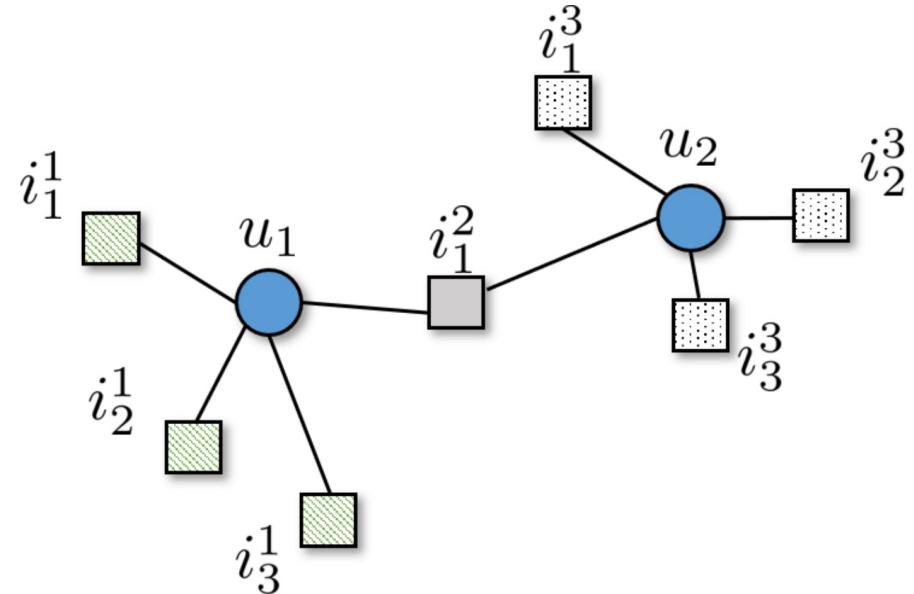
Area 3) Recommender systems

Example 3: Pinterest (social platform)

Image & User interaction in Pinterest



User-item interaction graph



Area 3) Recommender systems

Example 4: Other industry usecases

Applied Research Track

CIKM '20, October 19–23, 2020, Virtual Event, Ireland

P-Companion: A Principled Framework for Diversified Complementary Product Recommendation

Junheng Hao^{1,*}, Tong Zhao², Jin Li³, Junzhong Zhang¹, Qiang Cai¹, Yuxin Chen¹, Alireza Ebrahimi¹, Junzhou Huang¹
¹University of Wyoming, ²Wolfram Alpha, Inc., ³Amazon [jhaoyz@uwyo.edu, jli@uwyo.edu, qcai@uwyo.edu, ychen@uwyo.edu, ae@uwyo.edu, jhuang@uwyo.edu]



Applied Data Science Track Paper

ABSTRACT

Complementary product recommendation (CPR), aiming at providing product suggestions that are often bought together, or joint demand, forms a pivotal component of e-commerce service, however, existing methods are far from optimal. Given one product, how to recommend its complementary products of different types is the key problem we tackle in this work. We first conduct an analysis to correct the inaccurate assumptions adopted by existing work to show that co-purchased products are not always complementary and further propose a new strategy to generate clean distant supervision labels for CPR modeling. Moreover, to bridge the gap from existing work that CPR does not only need relevance modeling but also requires diversity to fulfill the whole purchase demand, we develop a deep learning framework, P-Companion, to explicitly model both relevance and diversity. More specifically, given one product with its product type, P-Companion first uses an encoder-decoder architecture to predict the user's purchase history and then performs a multi-task learning process to predict the user's purchase intent and the user's purchase history. Finally, P-Companion performs a recommendation process based on the user's purchase history and the user's purchase intent. The experimental results show that P-Companion significantly outperforms state-of-the-art methods.

Junheng Hao^{1,*}, Tong Zhao², Jin Li³, Junzhong Zhang¹, Qiang Cai¹, Yuxin Chen¹, Alireza Ebrahimi¹, Junzhou Huang¹
¹University of Wyoming, ²Wolfram Alpha, Inc., ³Amazon [jhaoyz@uwyo.edu, jli@uwyo.edu, qcai@uwyo.edu, ychen@uwyo.edu, ae@uwyo.edu, jhuang@uwyo.edu]

ConSTGAT: Contextual Spatial-Temporal Graph Attention Network for Travel Time Estimation

Xiaomin Fang, Jizhou Huang*, Fan Wang, Ling Liu, Baide Li, Chao Tang
{fangxiaomin01, huangjizhou01, wangfan01, liubaidi, tangchao}@ust.hk

ABSTRACT

The task of travel time estimation (TTE) is vital for route planning and traffic management. It is a challenging task to estimate the travel time for a given route and departure time, particularly in urban areas where traffic conditions are highly dynamic. In intelligent transportation systems such as GPS, traffic prediction, route planning, and ride-hailing services. This task is challenging because of many essential aspects, such as traffic prediction and contextual information. First, the accuracy of traffic prediction is strongly correlated with the traffic speed of the road segments in a route. Existing work mainly adopts spatial-temporal graph neural networks to improve the accuracy of traffic prediction, where spatial and temporal information is used separately. However, one drawback is that the spatial and temporal correlations are not fully exploited to obtain better accuracy. Second, contextual information of a route, i.e., the connections of adjacent road segments in the route, is an essential factor that impacts the driving speed. Previous work mainly uses sequential encoding models to address this issue. However, it is difficult to scale up sequential models to large-scale real-world datasets.



na
port
tion
You
sent
or n
are
and
I
nen
dict
syst
and
eme
t and
ABSTRACT
With the rapid growth of social media applications (Apps) in recent years, user engagement has become increasingly important. In this paper, we propose a flexible definition of user engagement based on future metric expectations. Next, we design an end-to-end neural framework, FATE, which incorporates three key factors that we identify to influence user engagement: namely friendships, user actions, and temporal dynamics. To achieve explainable engagement predictions, FATE is based on a tensor-based graph neural network (GNN), LSTM and a mixture attention mechanism, which allows for (1) predictive explanations and (2) causal explanations. We evaluate our model on two real-world datasets, and the results show that FATE outperforms state-of-the-art models. Finally, we conduct a user study to understand the return of existing users using different metrics, such as churn rate prediction [38] and lifespan analysis [39]. Others model user engagement with macroscopic features (e.g., demographic information) [1] and microscopic static features (e.g., user activities) [19]. Recently, Liu et al. [20] propose using dynamic action graphs, where nodes are in-App actions, and edges are transitions between actions, to predict future activity using a neural model.

Despite some success, existing methods generally suffer from the following: (1) They fail to model friendship dependencies or ignore user-user interactions when modeling user engagement. As users are connected in social Apps, their engagement affects each other [32]. For example, active users may keep posting new contents, which attract his/her friends and elevate their engagement. Thus, it is essential to capture friendship dependencies and user interactions when modeling user engagement. (2) Engagement objectives may differ across Apps and even across features. For example, an

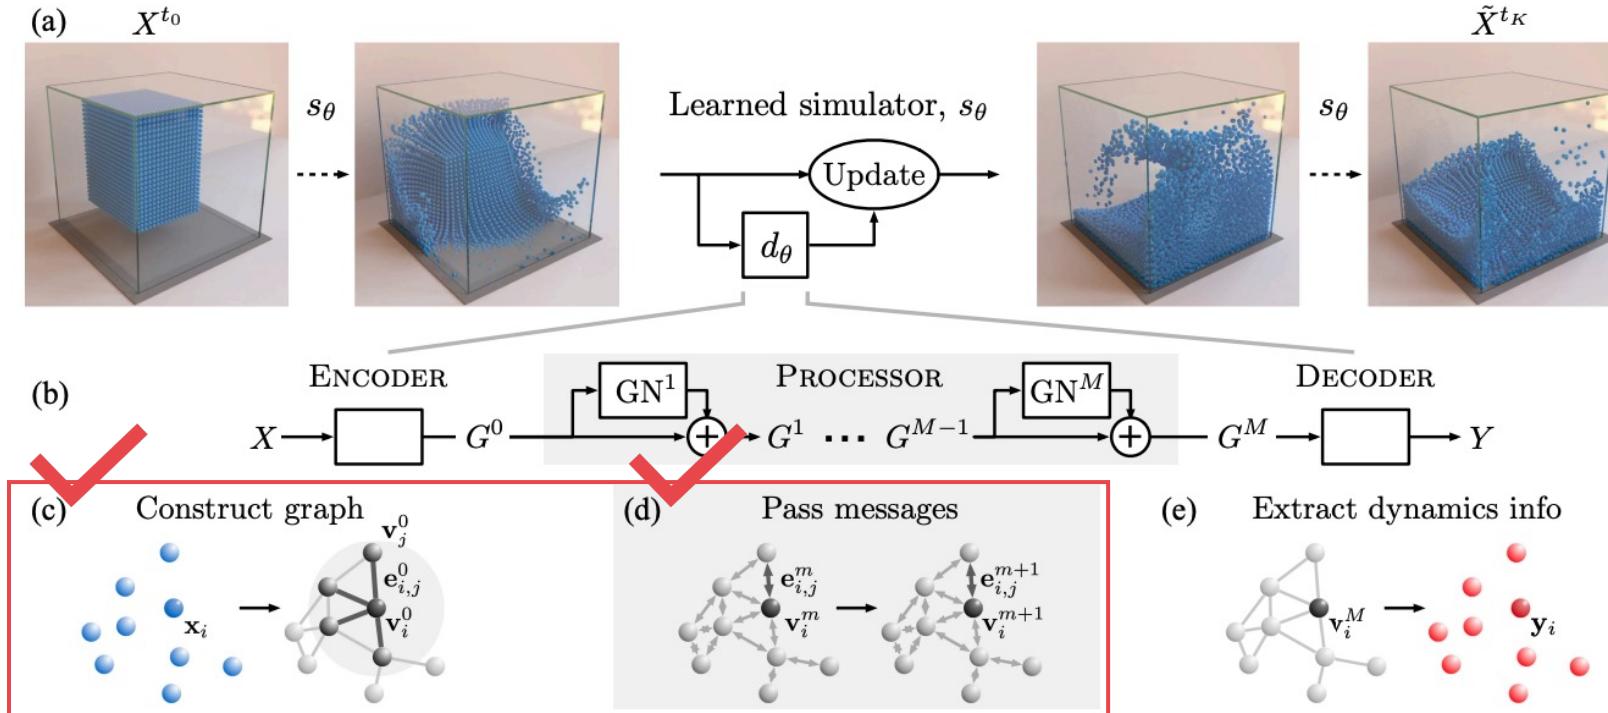
Xianfeng Zhou¹, Neil Shah², Xiaolin Shi³, Prasenjit Mitra¹, Suhang Wang^{4*}
Pennsylvania State University¹, Snp Inc.²
szwu14@psu.edu, fshah, xiaolin, mitra@snap.com
 **SnapChat**

Knowing your FATE: Friendship, Action and Temporal Explanations for User Engagement Prediction on Social Apps

15 Jun 2020

Area 4) Modeling physical systems

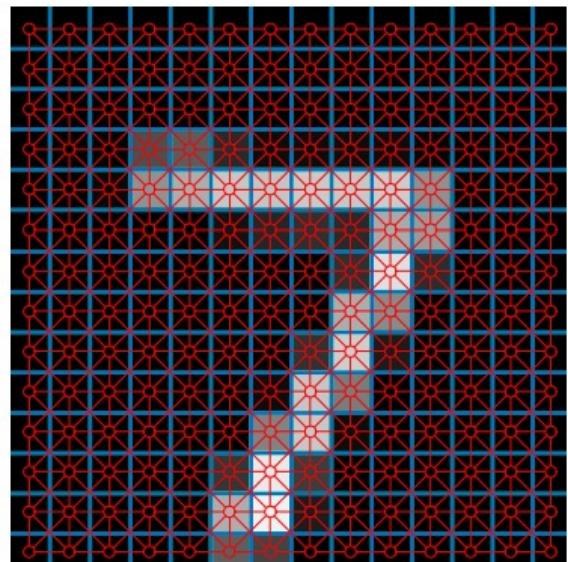
Example 5: Simulation of complex physical systems



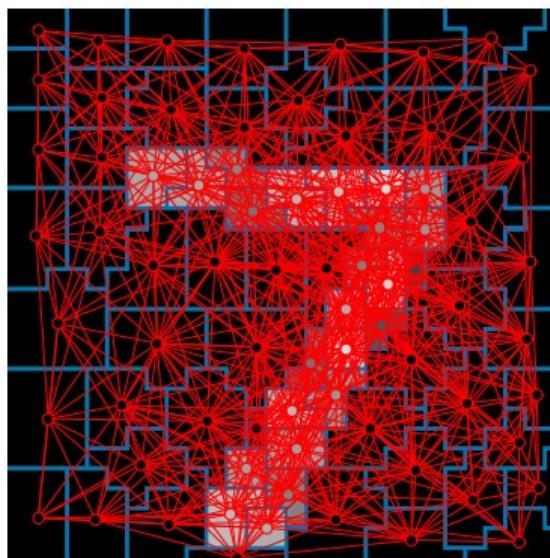
Similar to the ETA prediction task, how to construct the edges between particles will highly impact the rest of the learning process.

Area 5) Images are actually grid-like graphs

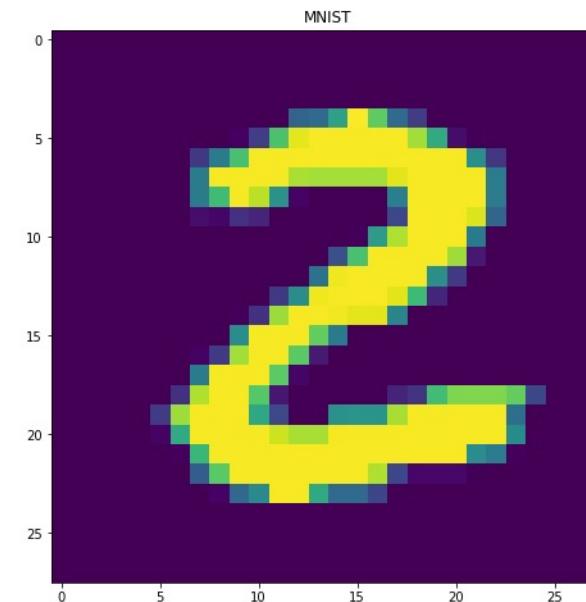
Example 6: MNIST and MNIST-sp



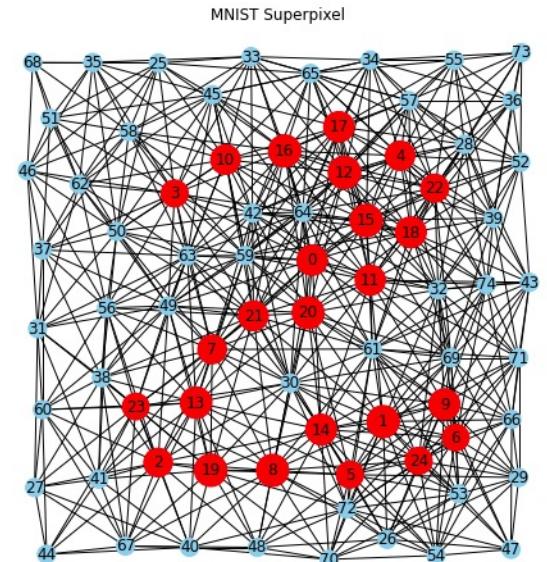
Regular grid



Superpixels



MNIST



MNIST Superpixel

MNIST-sp is quite commonly used as a benchmark dataset in the graph domain.

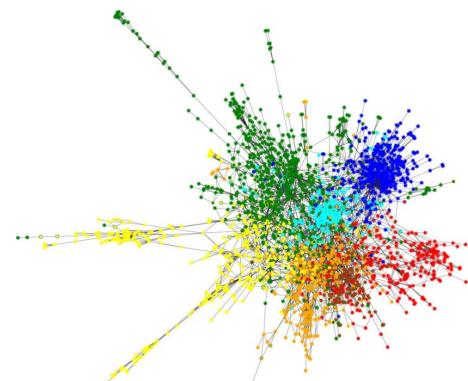
In academia: Benchmark datasets in the literature

Social



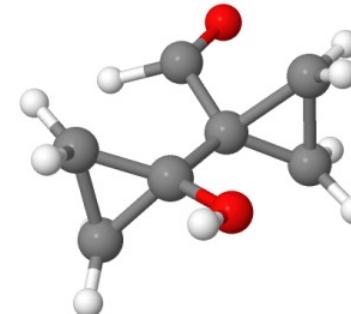
Node: People / Account
Edge: Connection
Node feature: Metadata

Citation / Web



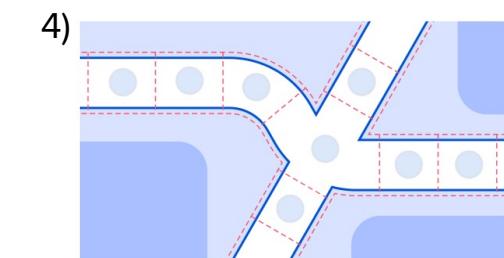
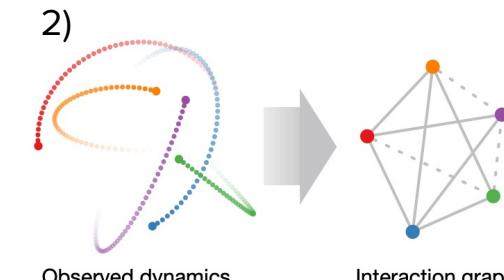
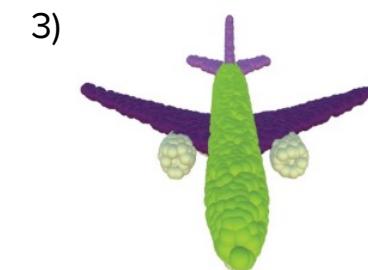
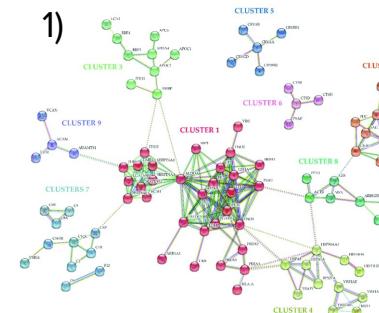
Node: Paper
Edge: Citation
Node feature: Abstract

Molecules



Node: Atom
Edge: Bond
Node feature: Atom type
Edge feature: Bond type

Biology / Simulation / etc.



Select benchmark datasets

- ***Planetoid dataset**
(Cora/Citeseer/Pubmed)
- **Coauthor**
- **WebKB**
(Texas/Cornell/etc.)
- **QM9**
- **Zinc**
- **MUTAG**

- 1) ****PPI** (protein-protein interaction)
- 2) Physical simulation (Kipf et al., 2018)
- 3) 3D point cloud (Wang et al., 2019)
- 4) Road network (Derrow-Pinion et al., 2021)

...and so much more

Yang et al., Revisiting Semi-Supervised Learning with Graph Embeddings, ICML 2016

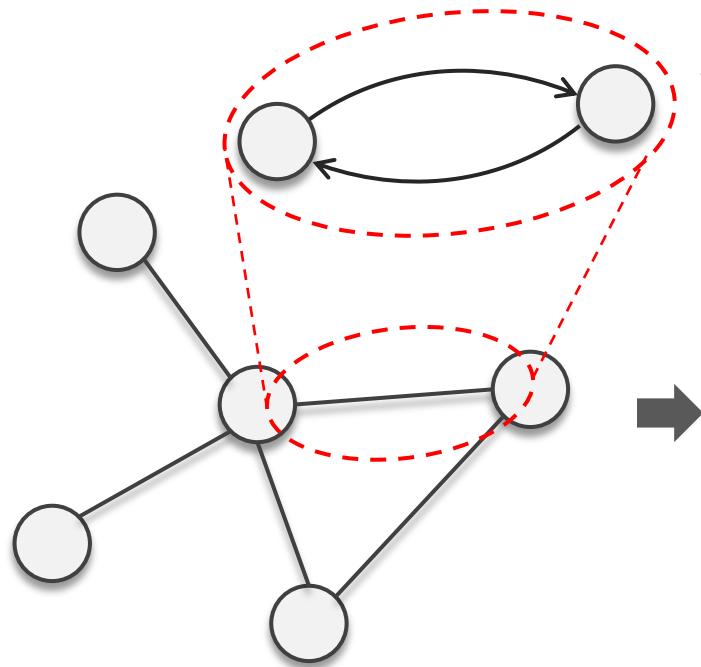
Kipf et al., Neural Relational Inference for Interacting Systems, ICML 2018

Wang et al., Dynamic Graph CNN for Learning on Point Clouds, ACM Transactions on Graphics 2019

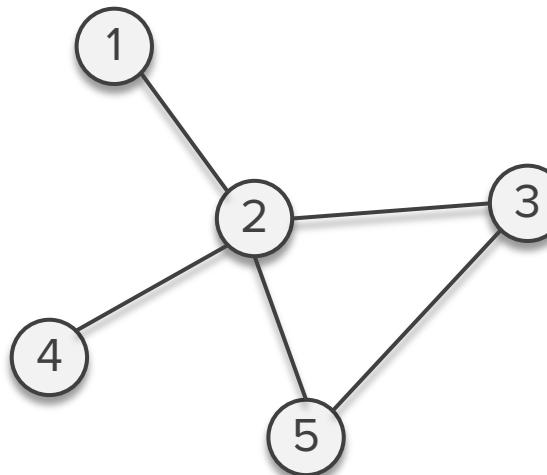
Derrow-Pinion et al., ETA Prediction with Graph Neural Networks in Google Maps, CIKM 2021

**Image source: https://www.researchgate.net/publication/324457787_iTRAQ_Quantitative_Proteomic_Analysis_of_Vitreous_from_Patients_with_Retinal_Detachment/figures?lo=1

Representing the graph as a adjacency matrix



*We treat undirected edges as
two directed edges going in both directions



	1	2	3	4	5
1	0	1	0	0	0
2	1	0	1	1	1
3	0	1	0	0	1
4	0	1	0	0	0
5	0	1	0	1	0

Undirected graph

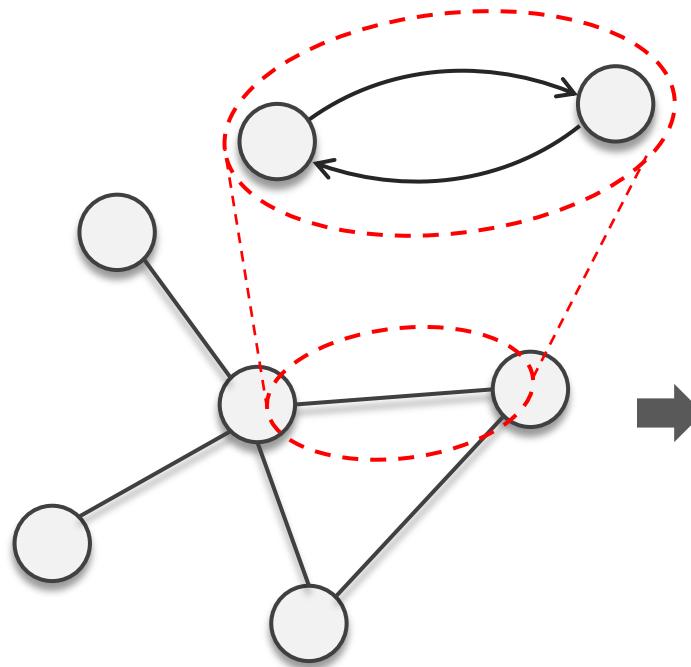
Assign arbitrary node ordering

- **Graphs with canonical node ordering is not common**
- Related research topic: Positional encoding in graphs
(Maskey et al., NeurIPS 2022)

Adjacency matrix

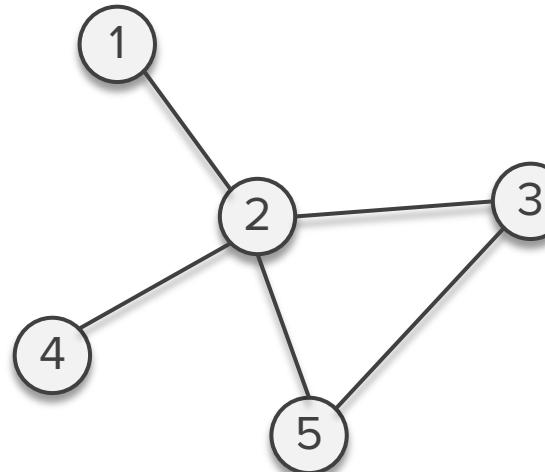
- Represent edge by assigning 1 for (i, j) -th element if node i and j are connected
- For weighted graphs: Assign a real number
- For graphs with multiple edges: Assign integers
- For directed graphs: Asymmetric matrix

Representing the graph as a adjacency matrix



Undirected graph

*We treat undirected edges as two directed edges going in both directions



$(1, 2), (2, 1), (2, 3), (3, 2), \dots$

Assign arbitrary node ordering

- **Graphs with canonical node ordering is not common**
- Related research topic: Positional encoding in graphs
(Maskey et al., NeurIPS 2022)

Edge list

- Represent graph by listing all edges
- Notice that for undirected edges, (i, j) and (j, i) both appear
- More **memory efficient** than (dense) adjacency matrix

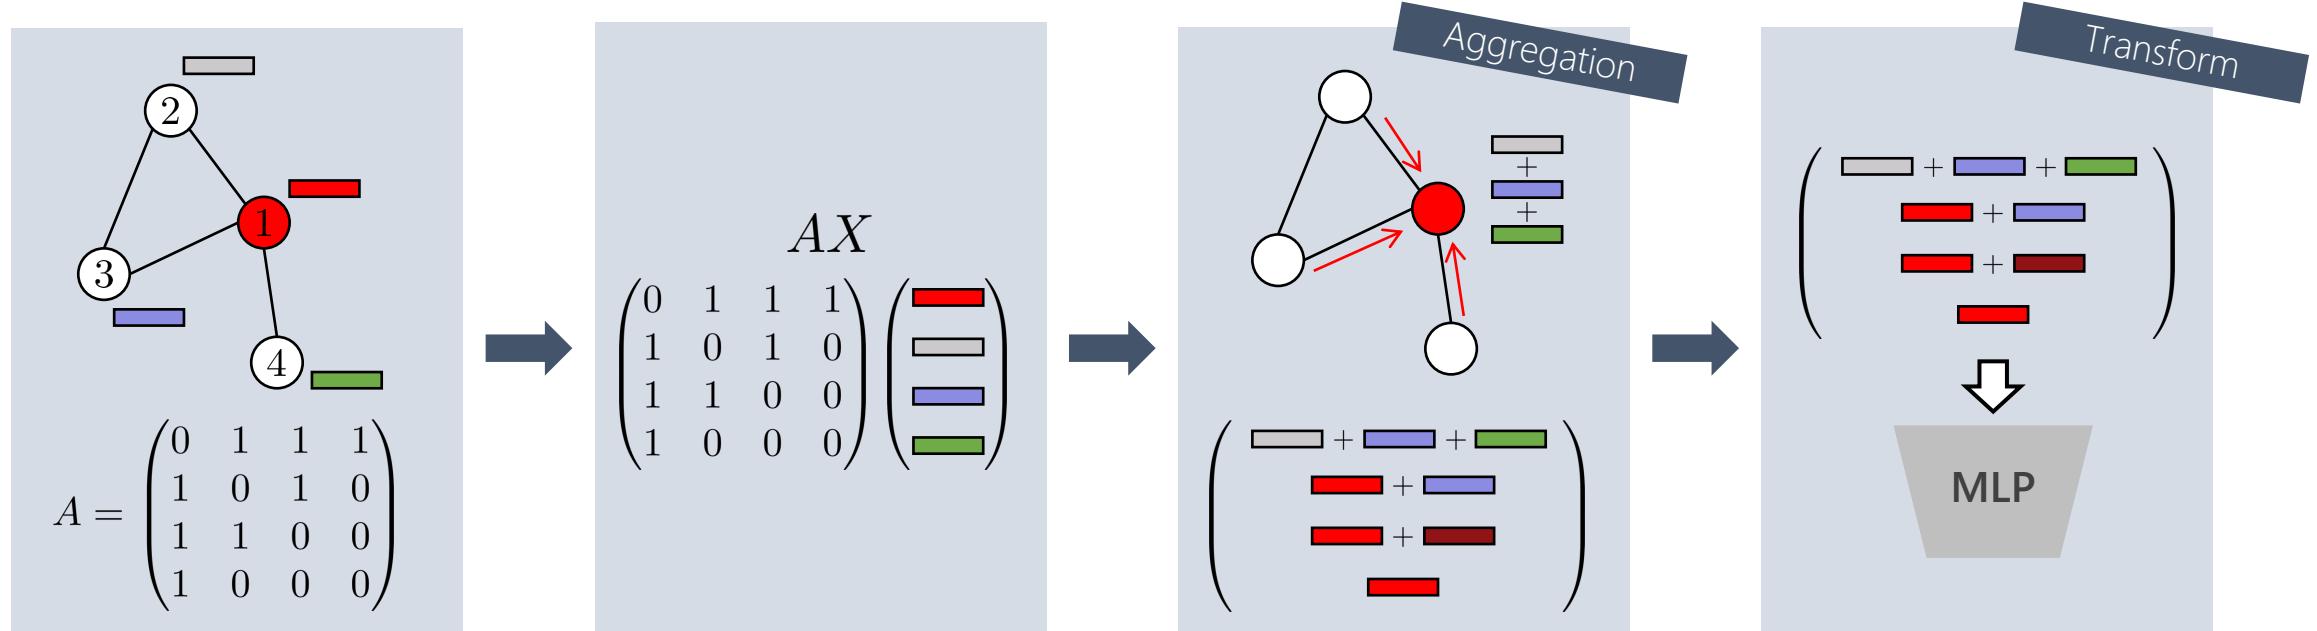
Part 1: A practical introduction to graphs and graph neural networks

Understanding of the general framework of graph neural networks (GNNs)

A simple, popular, and straightforward GNN

GCN (Graph Convolutional Network): Kipf & Welling, ICLR 2017

We are now ready to understand the basic principles of GNN, by looking at the most popular architecture.



Notice that, this whole procedure can be neatly expressed as: $\sigma(AX\Theta)$

Non-linear activation function $\sigma(\cdot)$

Adjacency matrix $A \in \mathbb{R}^{n \times n}$

Node feature matrix $X \in \mathbb{R}^{n \times d}$

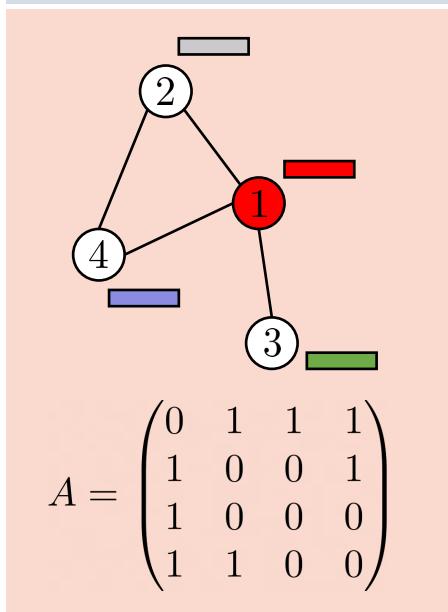
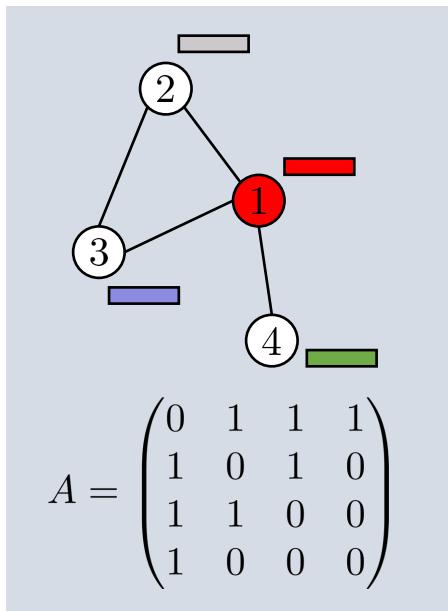
Learnable matrix $\Theta \in \mathbb{R}^{d \times d'}$

n : # of nodes

d : node feature dimensions

d' : dimension for the next layer

A deeper look into the node ordering problem

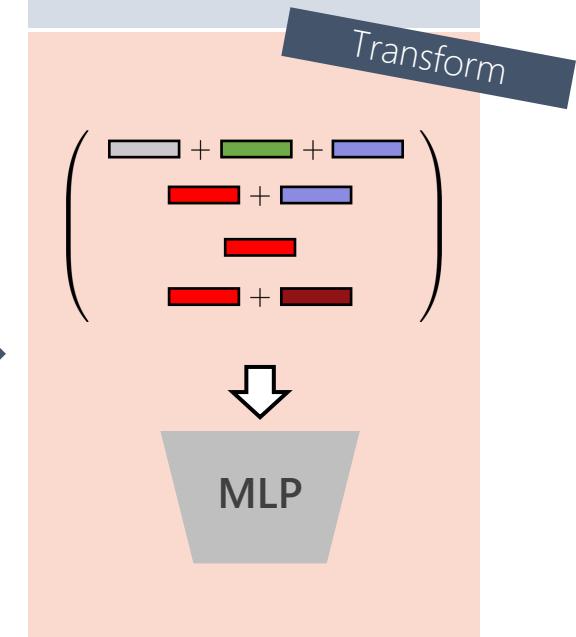
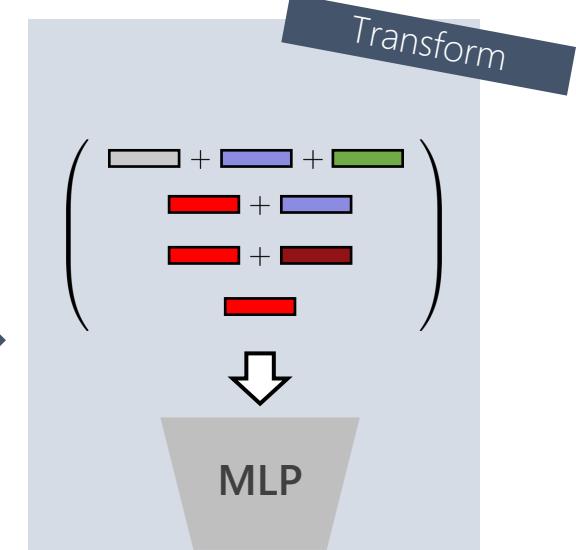
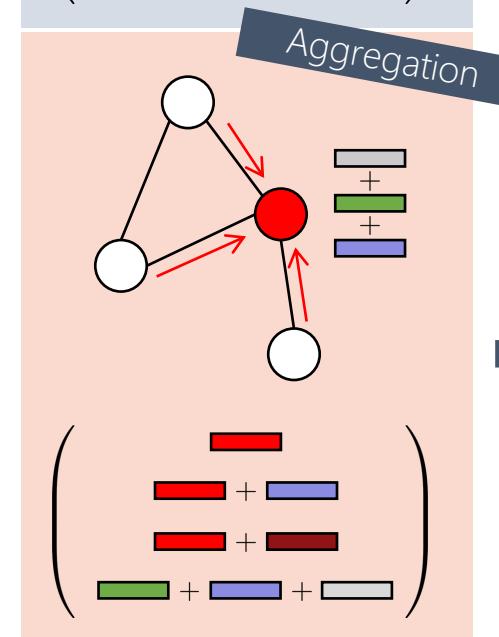
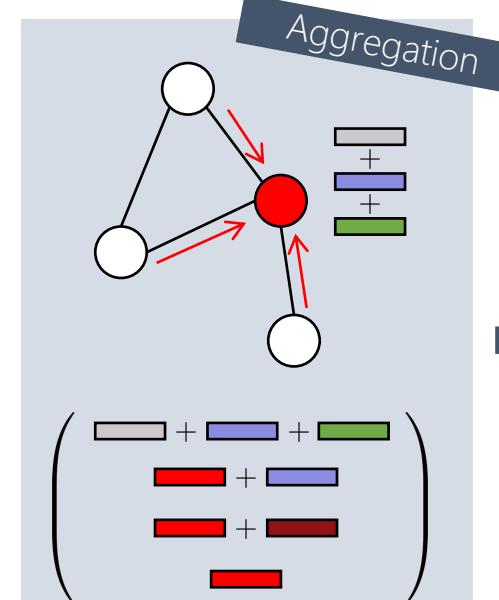


$$AX$$

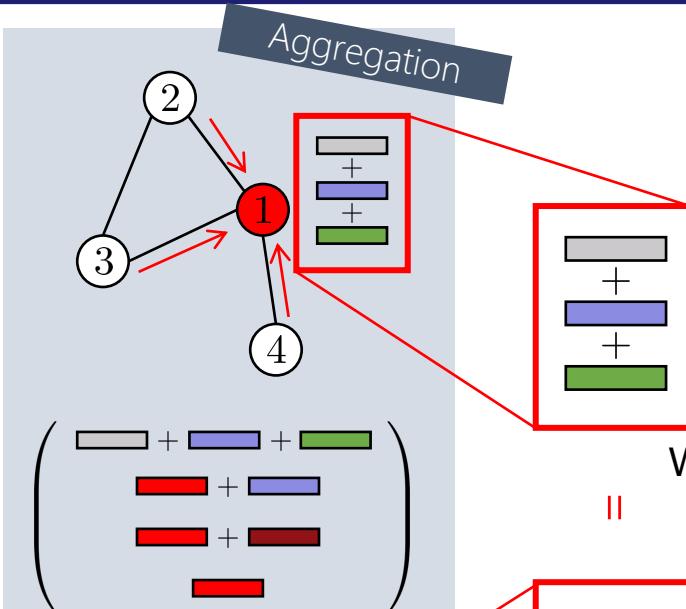
$$\begin{pmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \end{pmatrix} \begin{pmatrix} \text{red} \\ \text{grey} \\ \text{blue} \\ \text{green} \end{pmatrix}$$

$$AX$$

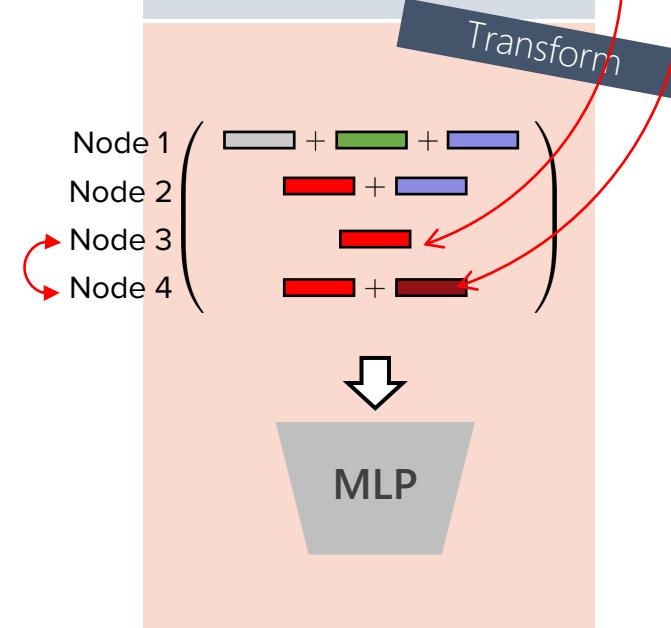
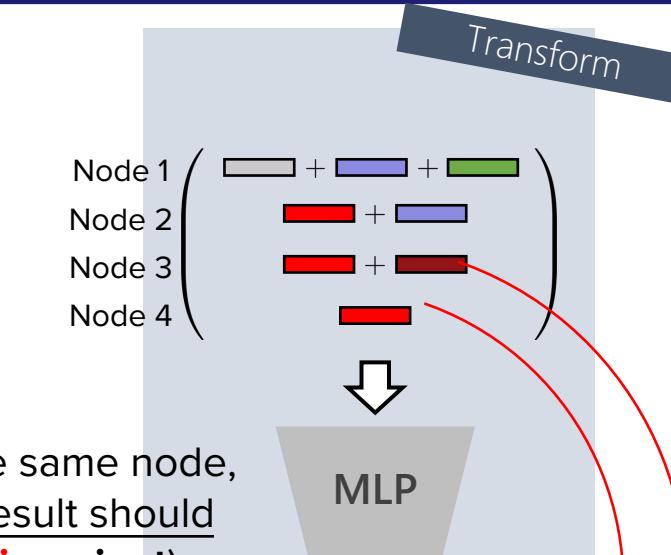
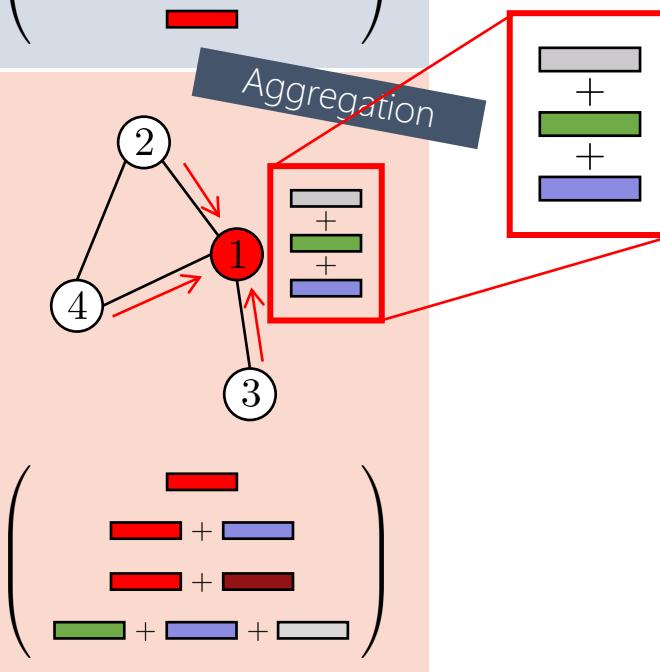
$$\begin{pmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 \end{pmatrix} \begin{pmatrix} \text{red} \\ \text{grey} \\ \text{blue} \\ \text{green} \end{pmatrix}$$



A deeper look into the node ordering problem



We are still computing for the same node,
therefore the aggregation result should
not change (permutation invariant)



As we have changed the node
order, this should also be
reflected in the embedding matrix
(permutation equivariant)

Practical design choices of GCN

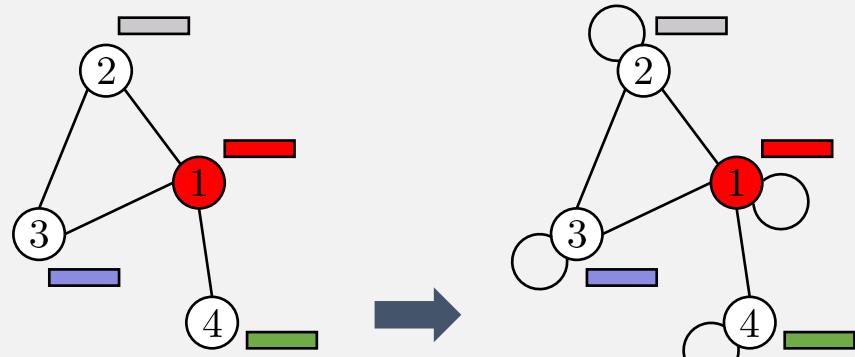
Of course, we can get creative with the graph structure to solve some practical issues

Problem 1: The information of the neighbor nodes are aggregated but not the node itself.

Problem 2: The adjacency matrix is not normalized, and the scale of the feature vectors may explode for repeated layers.

Resolution to problem 1

Add self-loops to each node

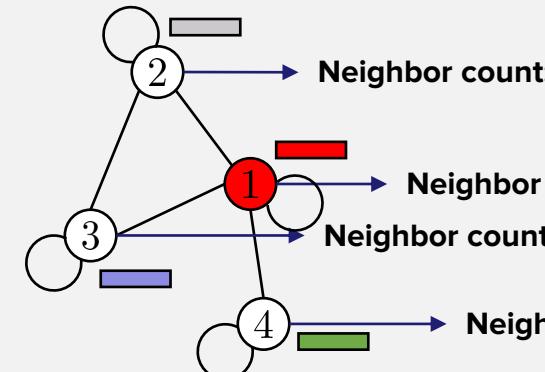


$$A = \begin{pmatrix} 0 & 1 & 1 & 1 \\ 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \end{pmatrix}$$

$$\hat{A} = A + I = \begin{pmatrix} 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 0 \\ 1 & 1 & 1 & 0 \\ 1 & 0 & 0 & 1 \end{pmatrix}$$

Resolution to problem 2

Normalization of \hat{A} :



$$\text{Neighbor count: } 3$$

$$\text{Neighbor count: } 4$$

$$\text{Neighbor count: } 3$$

$$\text{Neighbor count: } 2$$

$$D = \begin{pmatrix} 4 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 \\ 0 & 0 & 3 & 0 \\ 0 & 0 & 0 & 2 \end{pmatrix}$$

$$\tilde{A} = \hat{D}^{-1/2} \hat{A} \hat{D}^{-1/2} = \begin{pmatrix} \frac{1}{4} & \frac{1}{\sqrt{12}} & \frac{1}{\sqrt{12}} & \frac{1}{\sqrt{8}} \\ \frac{1}{\sqrt{12}} & \frac{1}{3} & \frac{1}{3} & 0 \\ \frac{1}{\sqrt{12}} & \frac{1}{3} & \frac{1}{3} & 0 \\ \frac{1}{\sqrt{8}} & 0 & 0 & \frac{1}{2} \end{pmatrix}$$

Final layer of GCN: $\sigma(\tilde{A}X\Theta)$

Abstraction: A general message-passing layer of GNNs

1. Message passing phase (Aggregation)

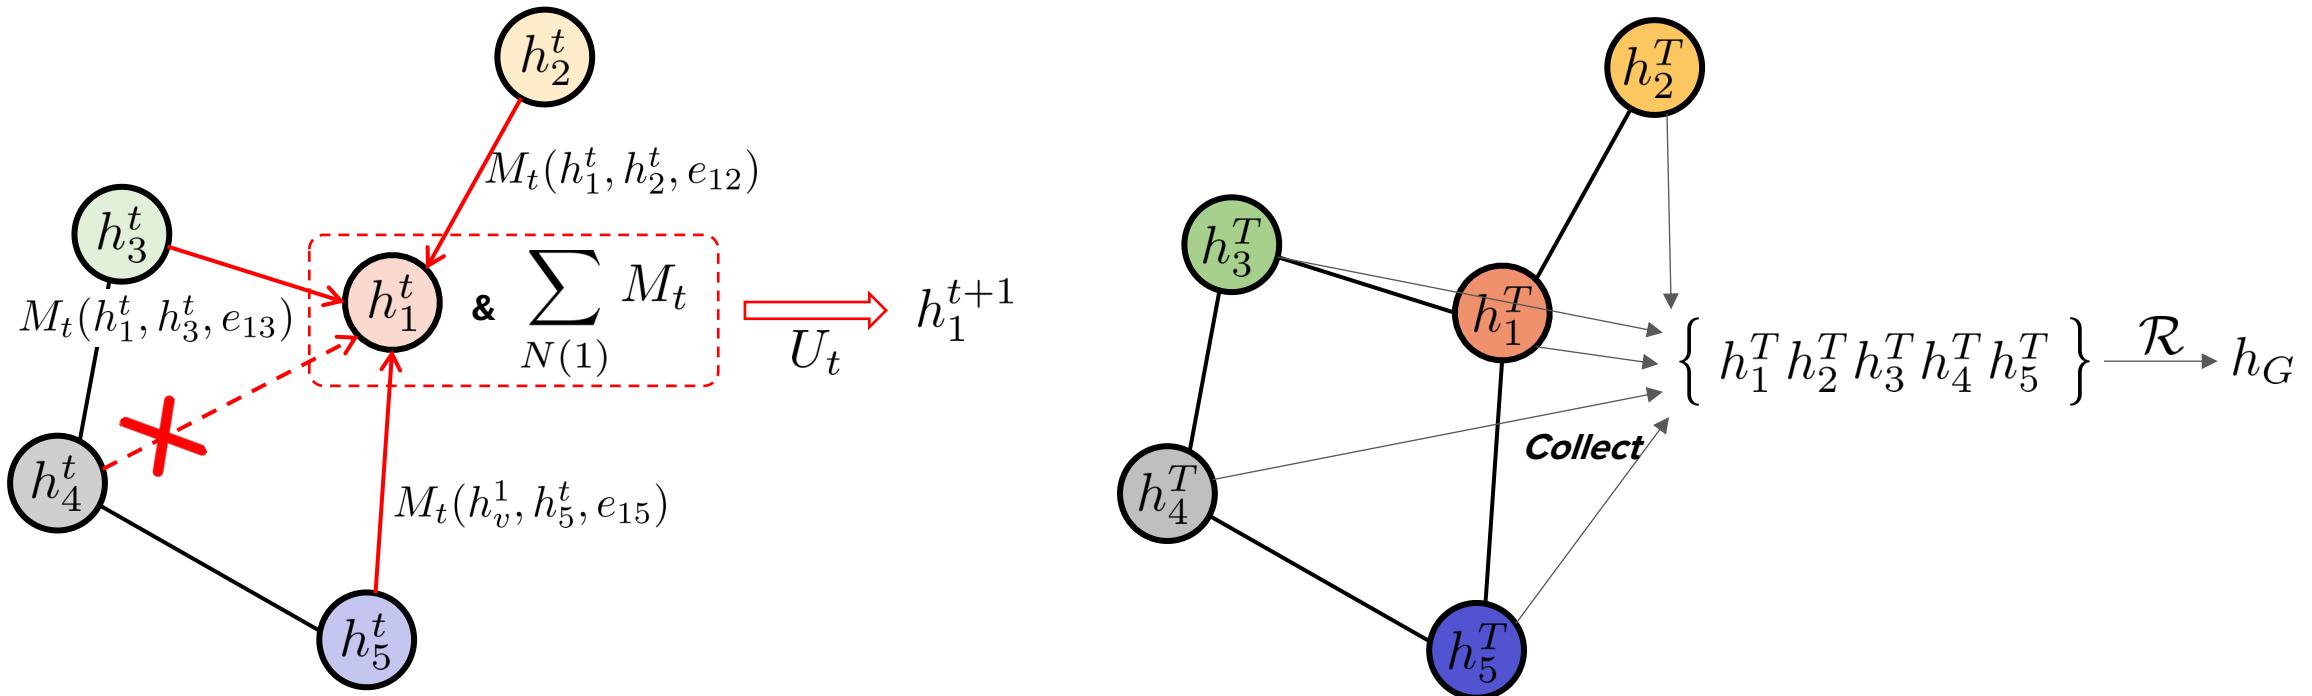
$$m_v^{t+1} = \sum_{w \in N(v)} M_t(h_v^t, h_w^t, e_{vw})$$

2. Update phase (Transformation)

$$h_v^{t+1} = U_t(h_v^t, m_v^{t+1})$$

3. Readout phase (Only for graph-level tasks)

$$h_G = \mathcal{R}(h_1^T, \dots, h_V^T)$$



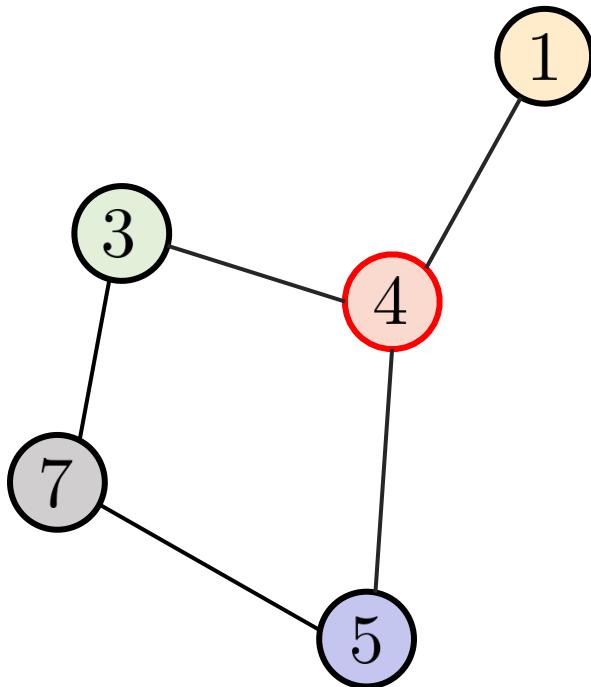
*Usually, we cite these papers for the term “message-passing”

[First formal introduction of the concept] Gilmer et al., “Neural Message Passing for Quantum Chemistry”, ICML 2017

[Comprehensive discussion & abstraction] Bronstein et al., Geometric Deep Learning: Grids, Groups, Graphs, Geodesics, and Gauges, arXiv 2021

Abstraction: A general message-passing layer of GNNs

GNN layer (Message-passing neural networks)



$$\mathbf{h}_u = \phi \left(\mathbf{x}_u, \bigoplus_{v \in \mathcal{N}_u} \psi(\mathbf{x}_u, \mathbf{x}_v) \right)$$



This operation must be permutation invariant to ensure the same result for different node orderings!

Summation / Average / Max pooling etc.

So if we re-describe GCN for node 4, it would be...

$$\mathcal{N}_u = \{1, 3, 5\} \cup \{4\} \quad \psi(\mathbf{x}_u, \mathbf{x}_1) = \frac{1}{\sqrt{2 \times 4}} \mathbf{x}_1 \quad \phi = \text{MLP}$$

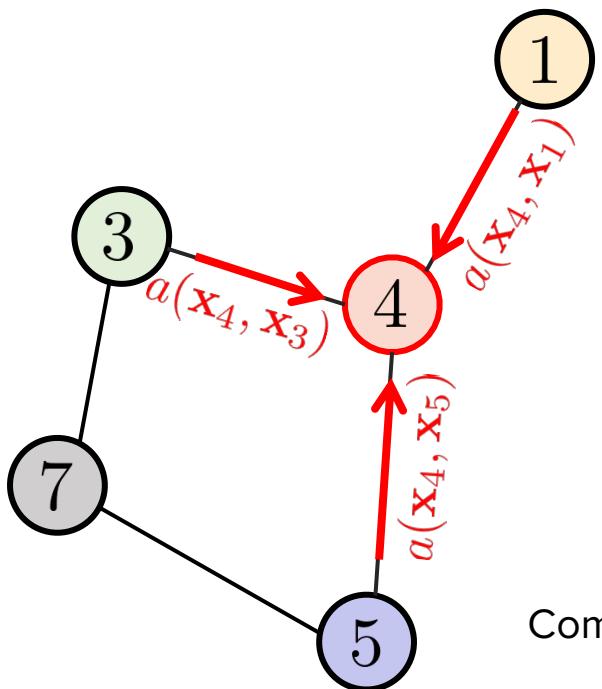
*Usually, we cite these papers for the term “message-passing”

[First formal introduction of the concept] Gilmer et al., “Neural Message Passing for Quantum Chemistry”, ICML 2017

[Comprehensive discussion & abstraction] Bronstein et al., Geometric Deep Learning: Grids, Groups, Graphs, Geodesics, and Gauges, arXiv 2021

Abstraction: A general message-passing layer of GNNs

Example: GAT (Veličković et al., ICLR 2018)



$$\mathcal{N}_u = \{1, 3, 5\} \cup \{4\}$$

Compare this part with GCN, the role of attention will be much more clear:

$$\mathbf{h}_u = \phi \left(\mathbf{x}_u, \bigoplus_{v \in \mathcal{N}_u} \boxed{\psi(\mathbf{x}_u, \mathbf{x}_v)} \right)$$

The model decides the strength of the ‘propagation’

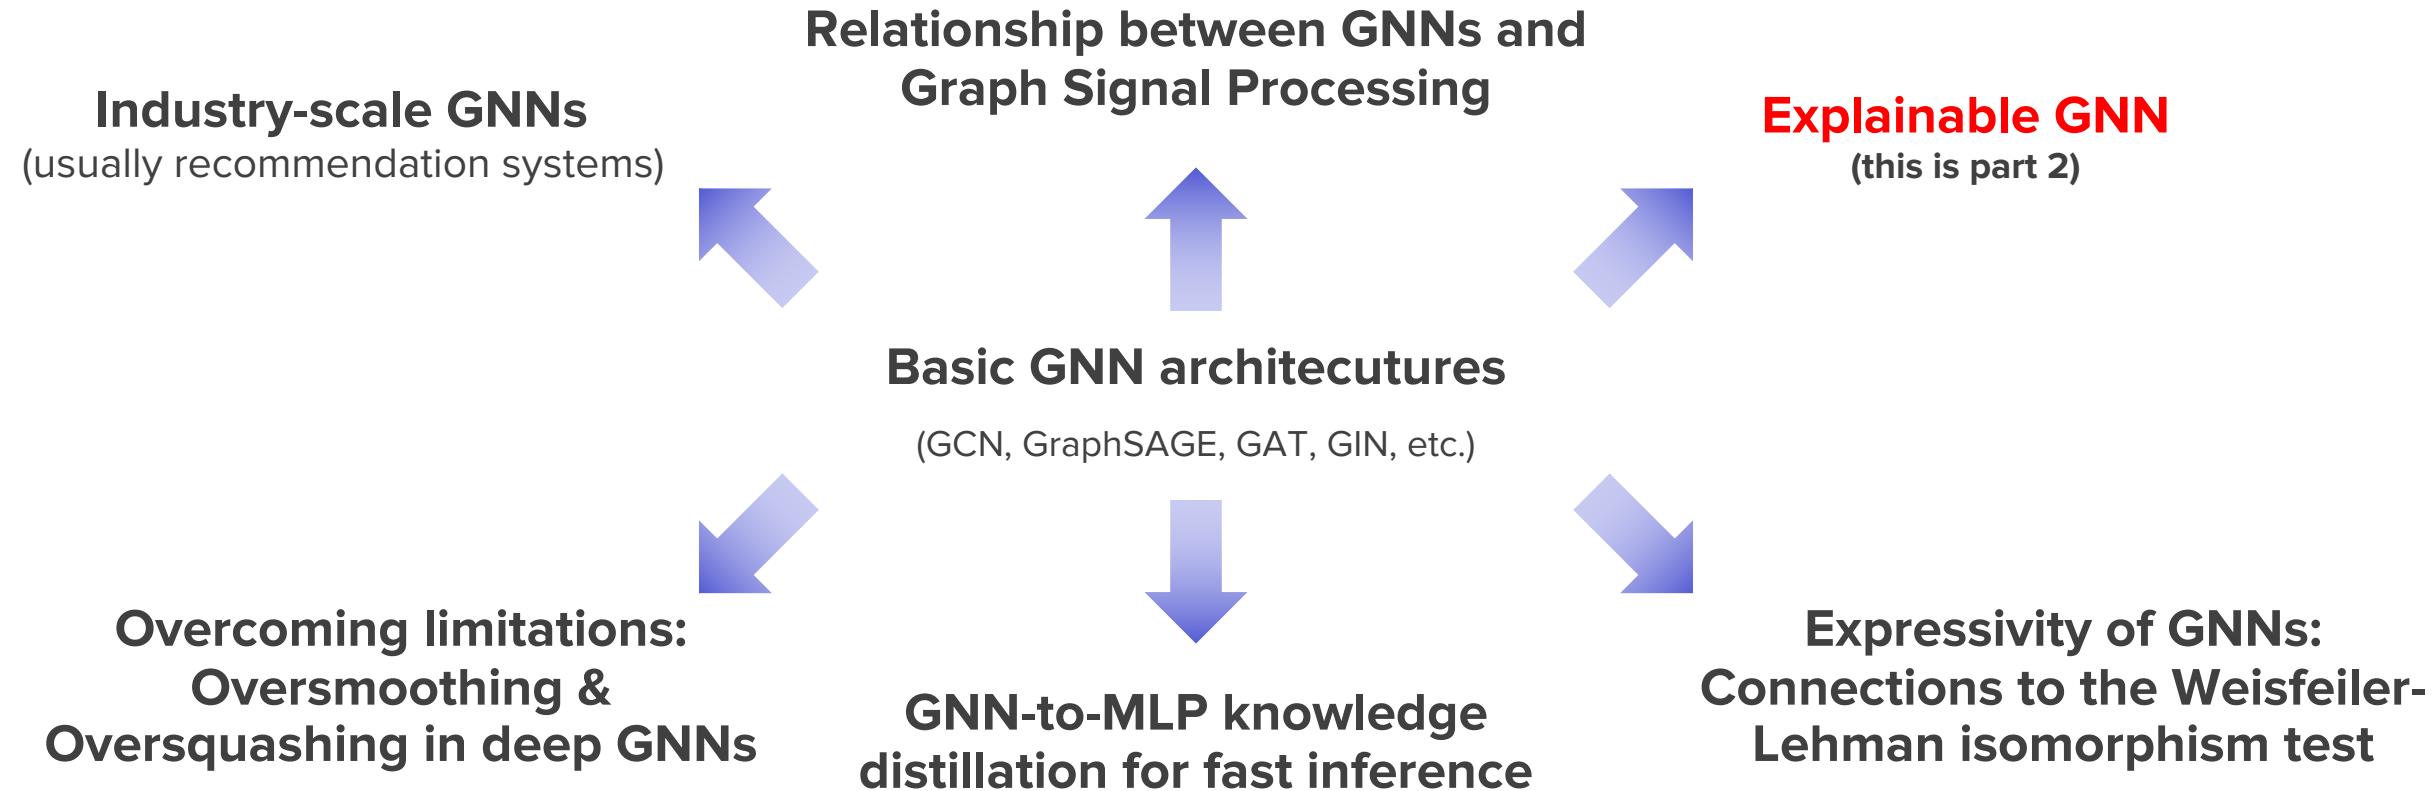
$$\psi(\mathbf{x}_u, \mathbf{x}_1) = \boxed{a(\mathbf{x}_u, \mathbf{x}_1)} \mathbf{x}_1 \quad \phi = \text{MLP}$$

$$\psi(\mathbf{x}_u, \mathbf{x}_1) = \frac{1}{\sqrt{2 \times 4}} \mathbf{x}_1$$

The node degree decides the strength of the ‘propagation’

There are a lot of fun & fundamental topics in the GNN literature

To name a few...



Before moving on, one slide on the library for graph learning

PyTorch Geometric ([link](#))

PyG Documentation

 PyG (PyTorch Geometric) is a library built upon  PyTorch to easily write and train Graph Neural Networks (GNNs) for a wide range of applications related to structured data.

It consists of various methods for deep learning on graphs and other irregular structures, also known as **geometric deep learning**, from a variety of published papers. In addition, it consists of easy-to-use mini-batch loaders for operating on many small and single giant graphs, **multi GPU-support**, `torch.compile` support, `DataPipe` support, a large number of common benchmark datasets (based on simple interfaces to create your own), the `GraphGym` experiment manager, and helpful transforms, both for learning on arbitrary graphs as well as on 3D meshes or point clouds.

- Jure Leskovec (Stanford/KumoAI/Snapchat)
- Faster library updates
- (Seems like) A larger community

Deep Graph Library ([link](#))



Framework Agnostic

Build your models with PyTorch, TensorFlow or Apache MXNet.



Efficient And Scalable

Fast and memory-efficient message passing primitives for training Graph Neural Networks. Scale to giant graphs via multi-GPU acceleration and distributed training infrastructure.

Diverse Ecosystem

DGL empowers a variety of domain-specific projects including `DGL-KE` for learning large-scale knowledge graph embeddings, `DGL-LifeSci` for bioinformatics and cheminformatics, and many others.

- Slower library updates
- Variable framework support
- Can be tricky to install older versions

- Additional library: NetworkX ([link](#)) – Library for **graphs in general**
 - Not a library for ML/DL
 - Often used in junction with PyG/DGL



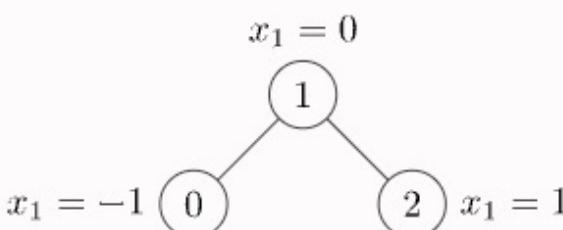
Wait, just one more slide on the library for graph learning

A very small PyG example

```
import torch
from torch_geometric.data import Data

edge_index = torch.tensor([[0, 1, 1, 2],
                          [1, 0, 2, 1]], dtype=torch.long)
x = torch.tensor([-1, 0, 1], dtype=torch.float)

data = Data(x=x, edge_index=edge_index)
>>> Data(edge_index=[2, 4], x=[3, 1])
```



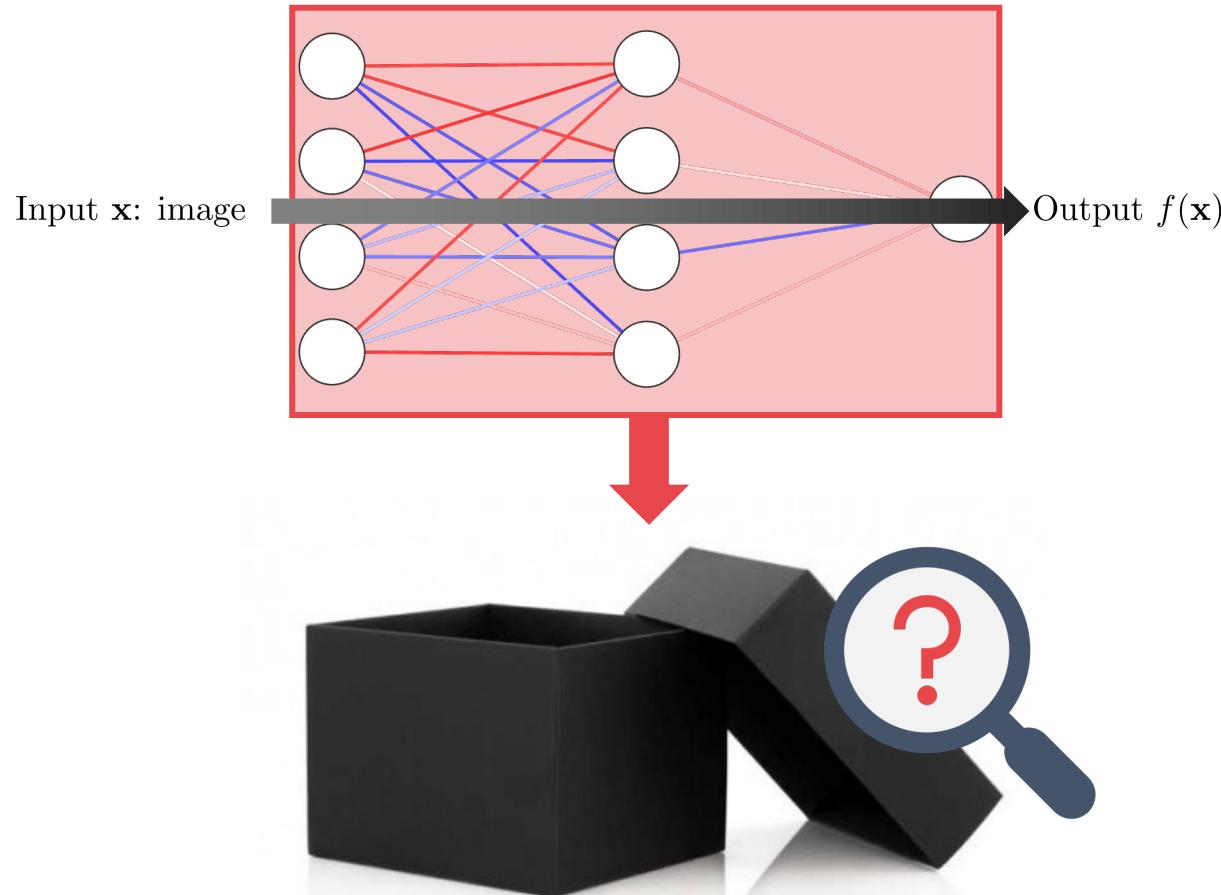
- You *at minimum* need to define `data.edge_index`
- Node features are usually represented as `data.x`
- Don't forget to include both directions for undirected graphs
- Most graph processing/manipulation tools are in `torch_geometric.utils`. Or just transform into a `networkx` object!

Part 2: Towards explainable graph learning with attention

Understanding the basic concepts of explainable AI

Why explainable AI?

Neural networks have complex structure with a lot of parameters.

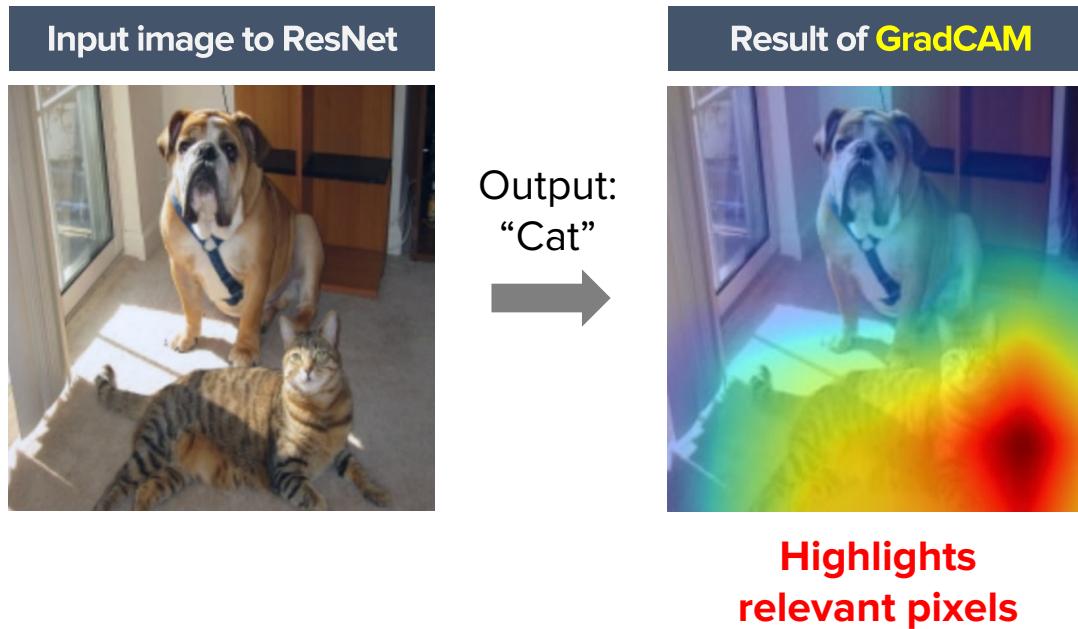


- Large number of parameters
 - High-nonlinearity
 - Complex inner structure
- has made them very ***hard to interpret and understand**, making it a **black box**.

Main question of Explainable AI:
Why has a neural network model made its prediction?

Attribution maps: The most popular type of explanation

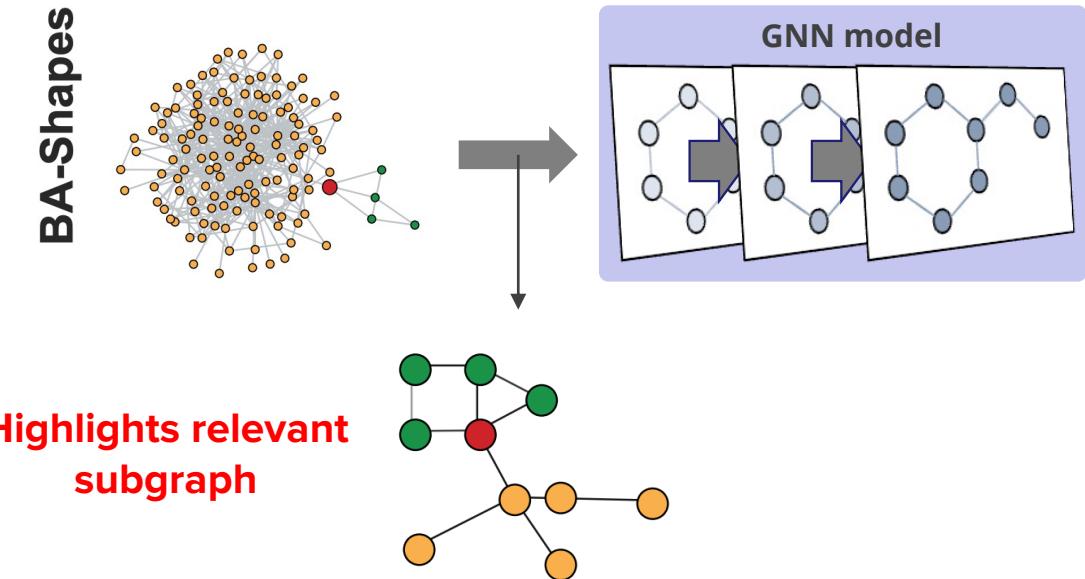
Attribution maps are one of the most popular ways, especially in CV and NLP.



Example: DTD [1], LRP [2], LIME [3], GradCAM [4], ...

Similar approaches are also popular in **GNN explanations**, too.

Computation graph



Example: GNNExplainer [5], PGExplainer [6], ...

[1] Montavon et al., "Explaining nonlinear classification decisions with deep Taylor decomposition", Pattern Recognit. 65: 211-22 (2017)

[2] Bach et al., "On Pixel-Wise Explanations for Non-Linear Classifier Decisions by Layer-Wise Relevance Propagation", PLOS ONE 10(7): e0130140.

[3] Ribeiro et al., "'Why Should I Trust You?': Explaining the Predictions of Any Classifier", KDD 2016

[4] Selvaraju et al., "Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization", ICCV 2017

[5] Ying et al., "GNNExplainer: Generating Explanations for Graph Neural Networks", NeurIPS 2019

[6] Luo et al., "Parameterized explainer for graph neural network", NeurIPS 2020

Part 2: Towards explainable graph learning with attention

Can we understand graph attention networks using attention?

What is attention?

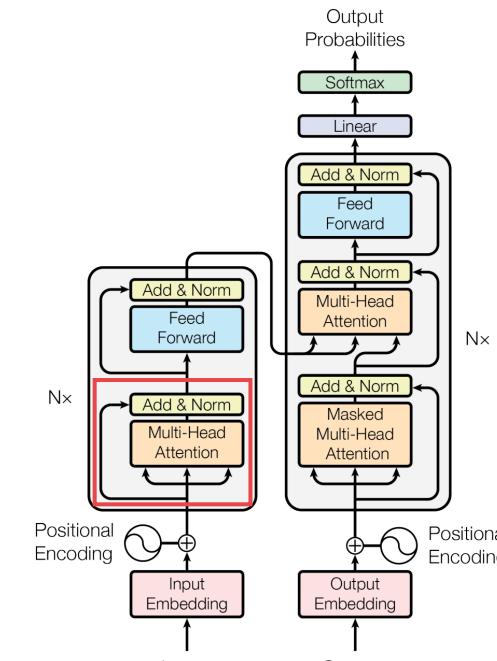
A weighted sum operation where the **weights** are determined by the model

The FBI is chasing a criminal on the run.
 The FBI is chasing a criminal on the run.
 The FBI is chasing a criminal on the run.
 The FBI is chasing a criminal on the run.
 The FBI is chasing a criminal on the run.
 The FBI is chasing a criminal on the run.
 The FBI is chasing a criminal on the run.
 The FBI is chasing a criminal on the run.
 The FBI is chasing a criminal on the run.
 The FBI is chasing a criminal on the run.

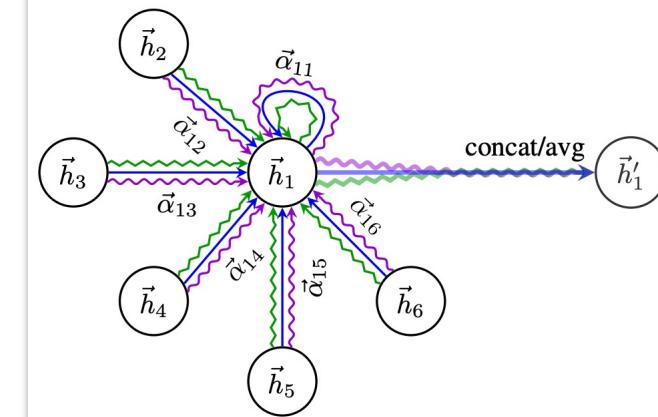
Figure 1: Illustration of our model while reading the sentence *The FBI is chasing a criminal on the run*. Color red represents the current word being fixated, blue represents memories. Shading indicates the degree of memory activation.

Long Short-Term Memory Networks (LSTMN)

(Cheng et al., EMNLP 2016)



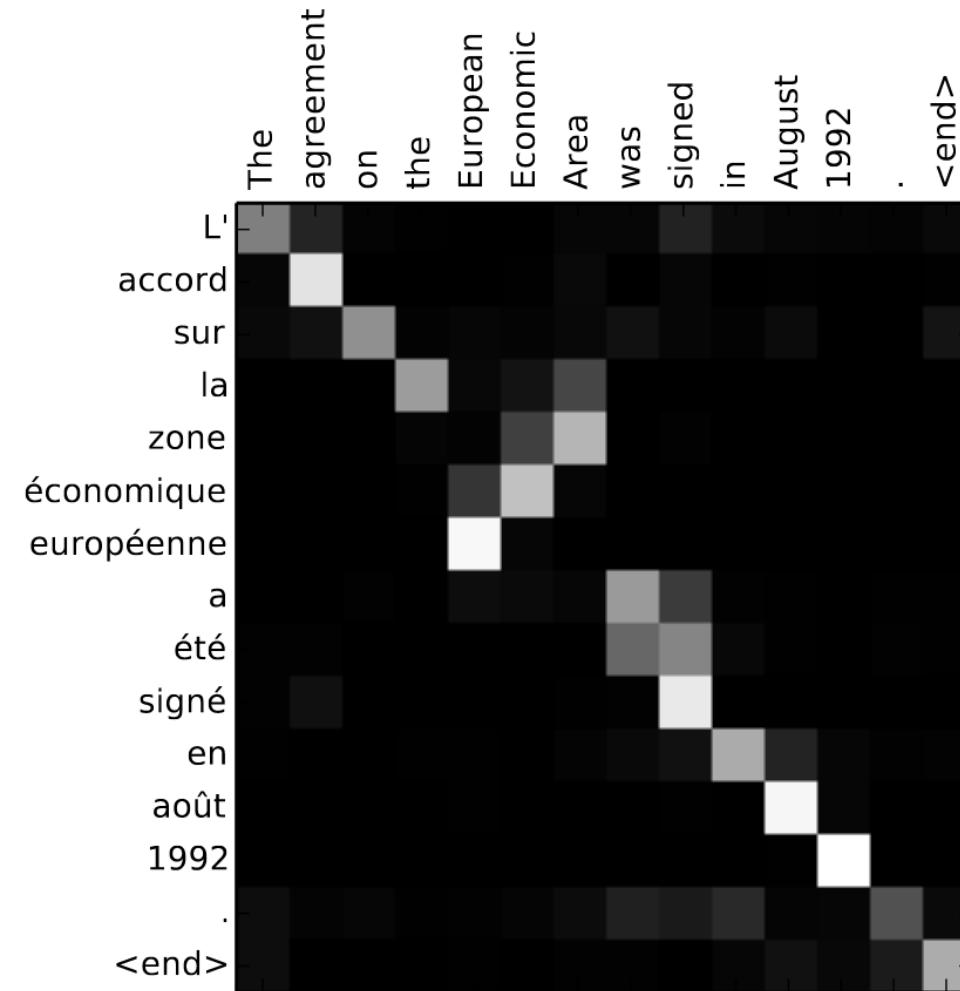
Transformer networks
(Vaswani et al., NeurIPS 2017)



Graph attention networks
(Velickovic et al., ICLR 2018)

Can we interpret attention = attribution?

**When we think of the role of attention, we can naturally interpret as 'where the model looks'
...which is essentially **attribution maps!** (at least intuitively)**



Can we just say attention = attribution?

Attention is heavily studied as an important candidate for model explanation

Is attention explanation?

Attention is not Explanation

Sarthak Jain
Northeastern University
jain.sar@husky.neu.edu

Byron C. Wallace

Attention is not not Explanation

Sarah Wiegreffe*

Yuval Pinter*

School of Interactive Computing
Georgia Institute of Technology
upv@gatech.edu

Is Attention Explanation? An Introduction to the Debate

Adrien Bibal, Rémi Cardon, David Alfter, Rodrigo Wilkens, Xiaou Wang,
Thomas François* and Patrick Watrin*
CENTAL, IL&C, University of Louvain, Belgium
{adrien.bibal, remi.cardon, thomas.françois, pat}@uclouvain.be

How Much Does Attention Actually Attend? Questioning the Importance of Attention in Pretrained Transformers

Michael Hassid[♡] Hao Peng^{◇*} Daniel Rotem[♡] Jungo Kasai[♣] Ivan Montero^{★*}
Noah A. Smith^{♦◇} Roy Schwartz[♡]
[♡]School of Computer Science & Engineering, Hebrew University of Jerusalem
[◇]Allen Institute for Artificial Intelligence [★]Apple, Inc.
[♣]Paul G. Allen School of Computer Science & Engineering, University of Washington
{michael.hassid, daniel.rotem, roy.schwartz}@mail.huji.ac.il
haop@allenai.org {jkasai, nasmith}@cs.washington.edu ivamon@apple.com

How to generate better attention heatmaps in transformers?

Quantifying Attention Flow in Transformers

Samira Abnar
ILLC, University of Amsterdam
s.abnar@uva.nl

Willem Zuidema
ILLC, University of Amsterdam

Transformer Interpretability Beyond Attention Visualization

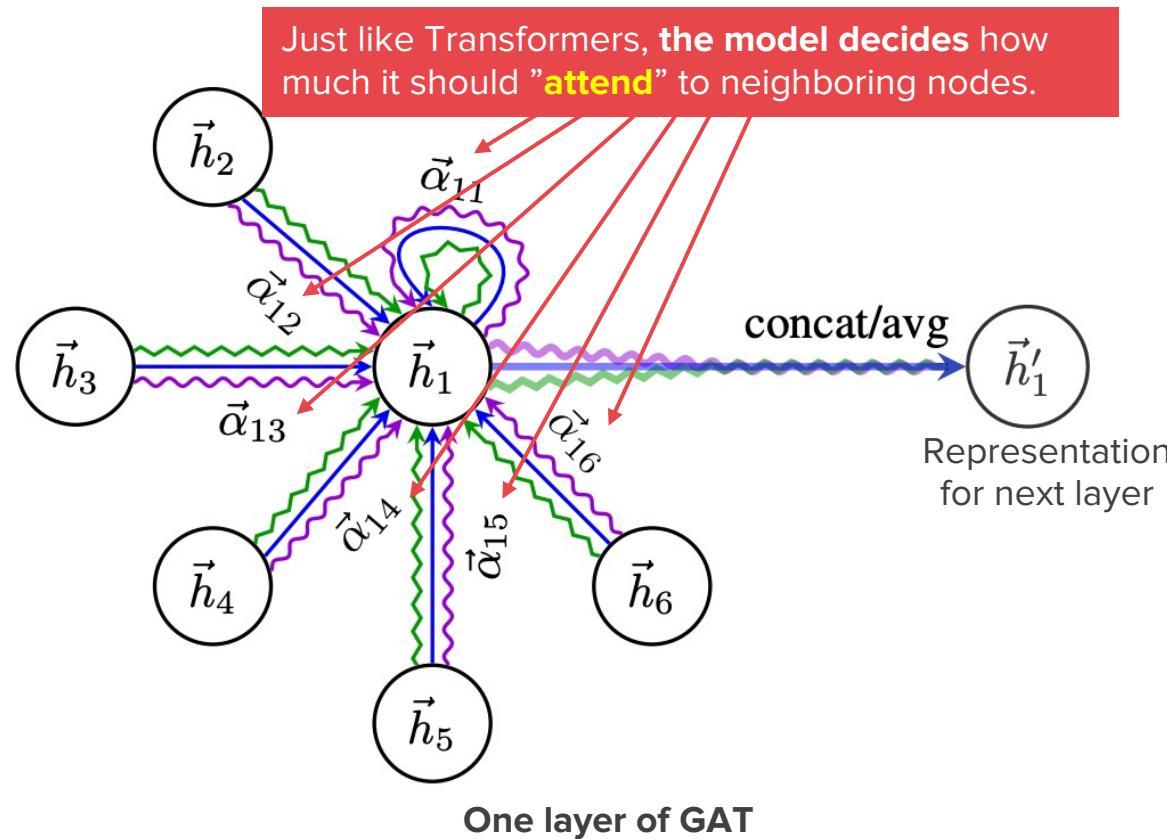
Hila Chefer¹ Shir Gur¹ Lior Wolf^{1,2}
¹The School of Computer Science, Tel Aviv University
²Facebook AI Research (FAIR)

Generic Attention-model Explainability for Interpreting Bi-Modal and Encoder-Decoder Transformers

Hila Chefer¹ Shir Gur¹ Lior Wolf^{1,2}
¹The School of Computer Science, Tel Aviv University
²Facebook AI Research (FAIR)

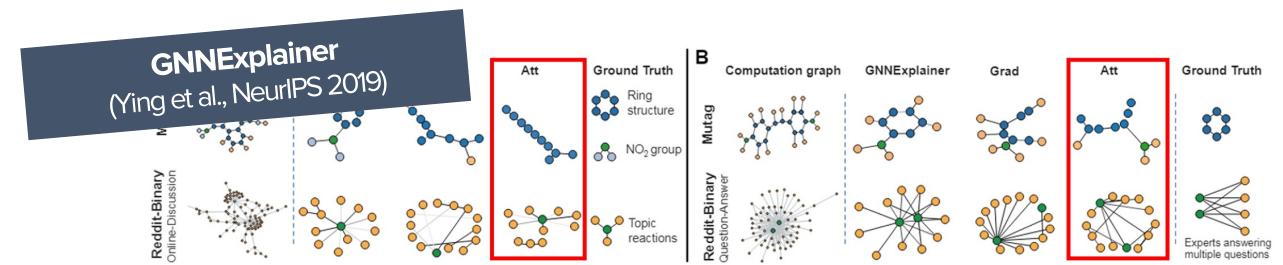
So graph attention networks (GATs) are explainable?

Well, the majority of the literature seems to overlook GATs as a valid candidate of ‘inherently explainable model’



GNN-XAI evaluation (Sanchez-Lengeling et al., NeurIPS 2020)							Node-level tasks								
	BA-Community			Tree-Grid			MPNN			GraphNets			GAT		
Baseline	0.27	0.27	0.27	GAT	0.27	0.27	0.38	0.38	0.38	0.38	0.62	0.62	0.62	0.62	
GradInput	0.58	0.64	0.39	GraphNets	0.72	0.52	0.51	0.5	0.5	0.5	0.65	0.71	0.66	0.67	
SmoothGrad(GI)	0.58	0.64	0.39	GAT	0.72	0.52	0.51	0.51	0.51	0.51	0.65	0.71	0.66	0.67	
GradCAM-last	0.79	0.84	0.86	0.8	0.7	0.67	0.68	0.68	0.68	0.68	0.7	0.77	0.81	0.7	
GradCAM-all	0.67	0.78	0.65	0.76	0.67	0.71	0.73	0.73	0.73	0.73	0.68	0.7	0.67	0.68	
IG	--	--	--	--	0.81	0.75	0.72	0.72	0.72	0.72	--	--	--	--	
Attention Weights	--	--	--	0.5	--	--	--	--	--	0.5	--	--	--	0.49	

“...have several blocks and attention heads, so for each component we take their average to combine them to a scalar value assigned to each edge.”



“...it is not obvious which attention weights need to be used for edge importance, Each edge’s importance is thus computed as the average attention weight across all layers.”

Explanation AUC						
PGEExplainer	0.750	0.905	0.612	0.717	0.783	
PGEExplainer	0.739	0.824	0.667	0.674	0.765	
PGEExplainer	-	-	-	0.773	0.653	
PGEExplainer	0.925	0.836	0.948	0.875	0.742	0.727
Improve	4.1%	13.0%	4.1%	3.7%	24.7%	11.5%
Inference Time (ms)						
GNNExplainer	650.60	696.61	690.13	713.40	934.72	409.98
PGEExplainer	10.92	24.07	6.36	6.72	80.13	9.68
Speed-up	59x	29x	108x	106x	12x	42x

“Each edge’s importance is obtained by averaging its attention weights across all attention layers.”

(Left) Velickovic et al., Graph attention networks, ICLR 2018

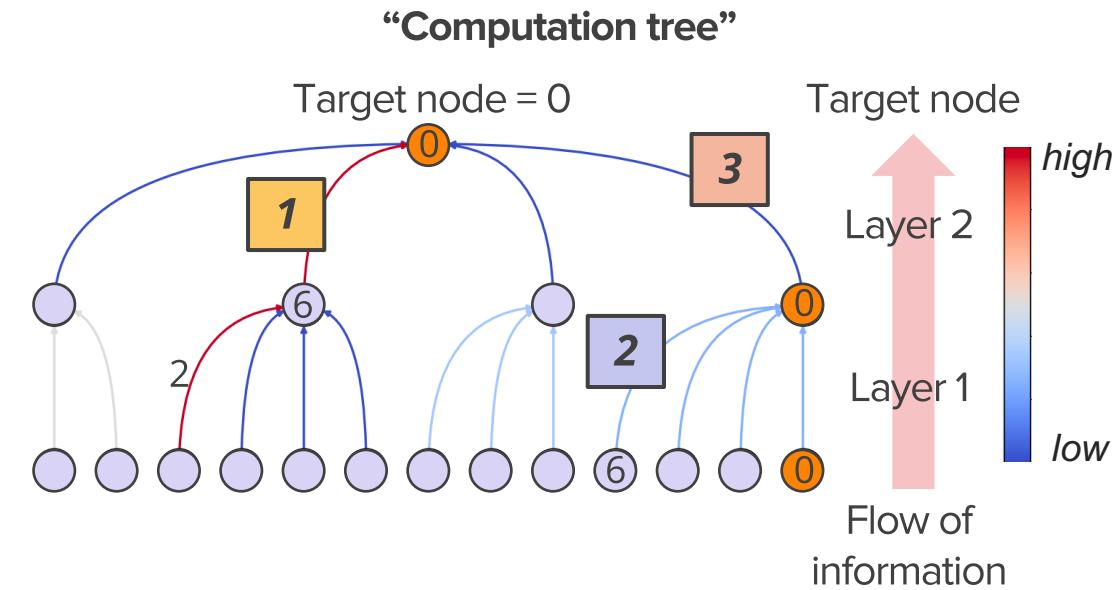
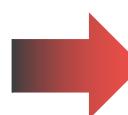
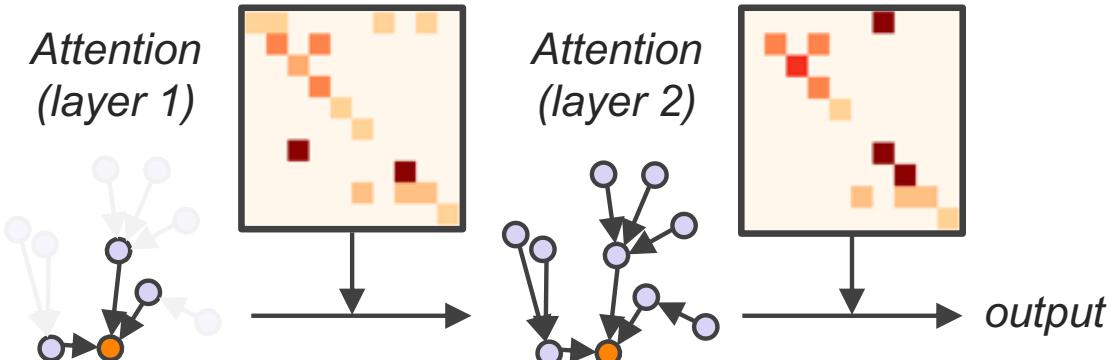
Sanchez-Lengeling et al., Evaluating attribution for graph neural networks, NeurIPS 2020

Ying et al., GNNExplainer: Generating explanations for graph neural networks, NeurIPS 2019

Luo et al., Parameterized Explainer for Graph Neural Network, NeurIPS 2020

GATs are explainable... with a little bit of extra effort

We just need to consider the ‘flow’ of information better within the GAT model



Proposed calculation (Shin et al., AAAI 2025)

$$\phi_{6,0}^0 = \boxed{1} + \boxed{3} \times \boxed{2}$$

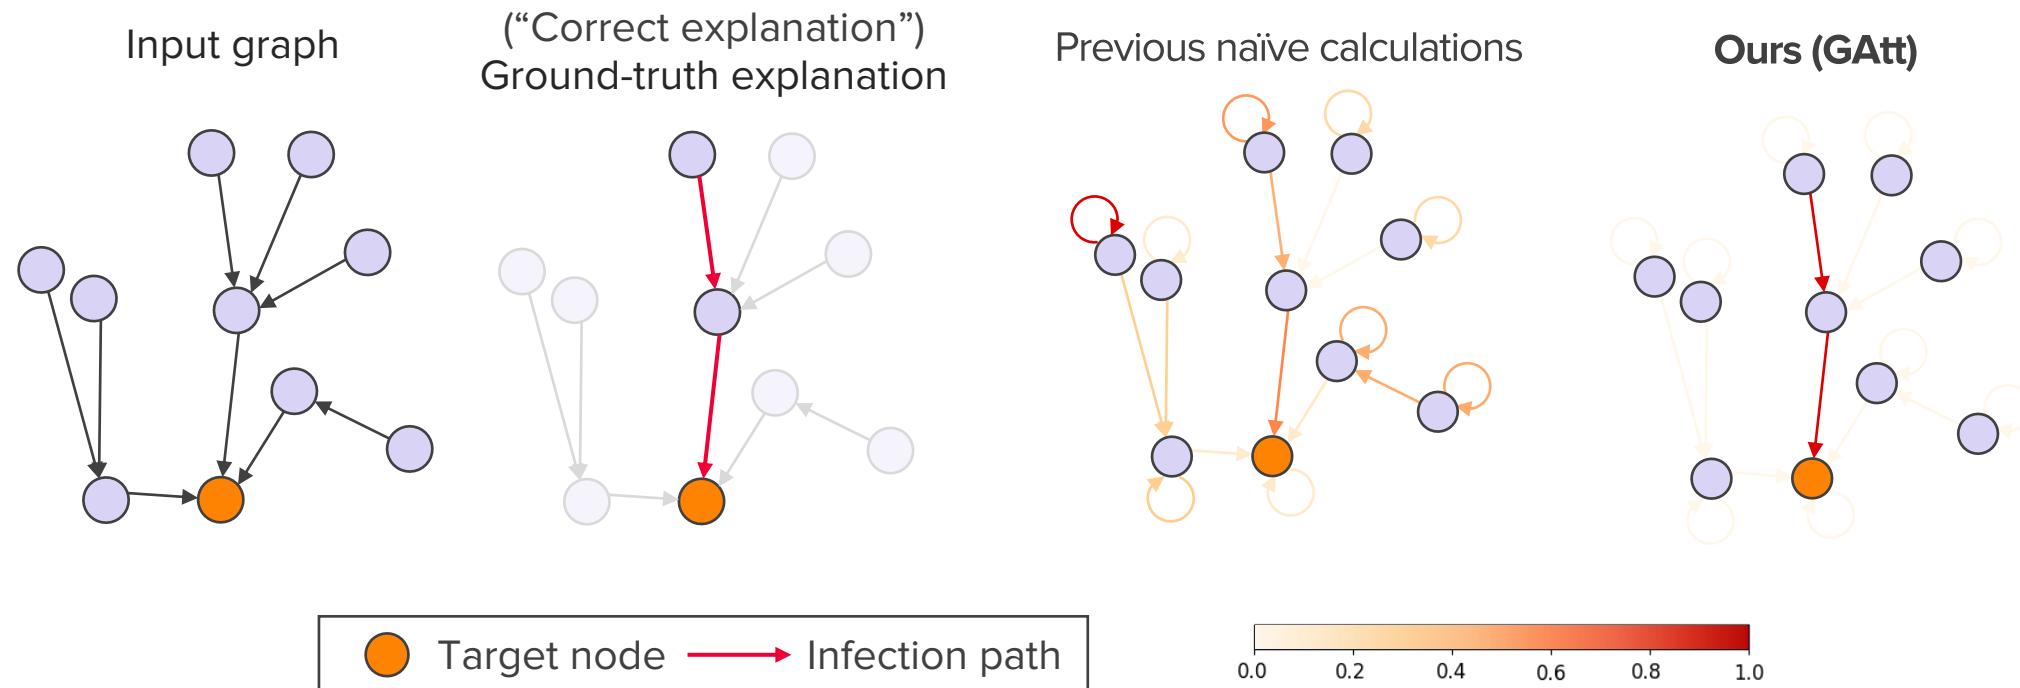
“Importance of edge (6, 0)
when the target node is 0”

1. Add all occurrences
in the computation tree

2. Multiply other attention weights
along the flow of information

And we can immediately get better attribution maps

Case study: Infection dataset (Faber et al., KDD 2021)



*Note: We do not change the GAT model. Remember, our contribution is how to calculate attribution maps after the training is complete

Takeaways

Part 1: A practical introduction to graphs and graph neural networks

1. Understanding of **graphs** as a **general data type**
 - Nodes & Edges (“connections”)
 - A lot of things can be represented as a graph, including images!
2. Understanding of the general framework of **graph neural networks (GNNs)**
 - Message-passing = GNN (Unless it’s a graph transformer)
 - Aggregation + Transformation

Part 2: Towards explainable graph learning with attention

1. Understanding the basic concepts of **explainable AI**
 - Attribution maps = “Important parts of the input”
 - A lot of GNN explanations are also attribution maps
2. **Answer to the question: Can we understand graph attention networks using attention?**
 - (Shin et al., AAAI 2025) Conclusion: YES, but with a little bit of effort
 - BTW, this conclusion is applicable to other GNNs with self-attention

Thank you!

Please feel free to ask any questions :)

jordan7186.github.io

jordan3414@yonsei.ac.kr